



# Module 9 Part 2

## NoSQL Databases

**BITS Pilani**

Harvinder S Jabbal  
SSZG653 Software Architectures

# NoSQL databases



- While **Relational databases** (SQL databases) are good for transaction management, not all applications need this feature.
- Many **web applications** such as marketing applications, IoT applications, **need fast processing** of data but do not need transaction management
- NoSQL databases came as a response to this need

<https://www.dataversity.net/a-brief-history-of-non-relational-databases/#>

November 9, 2024

SE Z06511-SS Z0653 Software  
Architectures

# NoSQL databases



Traditional databases also have limitations on storage

NoSQL databases are

- Scalable (Sharding)
- Fast (In Memory)
- Available (Replication)
- Handle semi-structured and unstructured data
- Rapidly adapt to changing data needs

However most lack

- Transaction support (ACID)- Atomicity, Consistency, Isolation, Durability.
- Join feature

# Examples of usage of NoSQL DB



- **Tesco**, Europe's #1 retailer, uses NoSQL to manage **millions of products**, promotions supply chain, etc.
- **Ryanair**, the world's busiest airline, uses NoSQL to **power its mobile app** serving over 3 million users
- **Marriott** is deploying NoSQL for its **reservation system** that books \$38 billion annually
- **Gannett** (USA Today) the #1 U.S. newspaper publisher, uses NoSQL for its proprietary **content management** system, Presto
- **GE** is deploying NoSQL for its Predix platform to help manage the **Industrial Internet**

# NoSQL - Flexibility



In RDBMS, if we want to add a new column, we need to change the schema

## Iteration 1 — First, Last

Schema Utilized

USERS		
ID	First	Last

Brendan	Bond
---------	------



USERS		
0	Brendan	Bond

## Iteration 2 — First, Last, *Twitter*

Schema Utilized

USERS		
ID	First	Last

Shane	Johnson	@shane_dev
-------	---------	------------



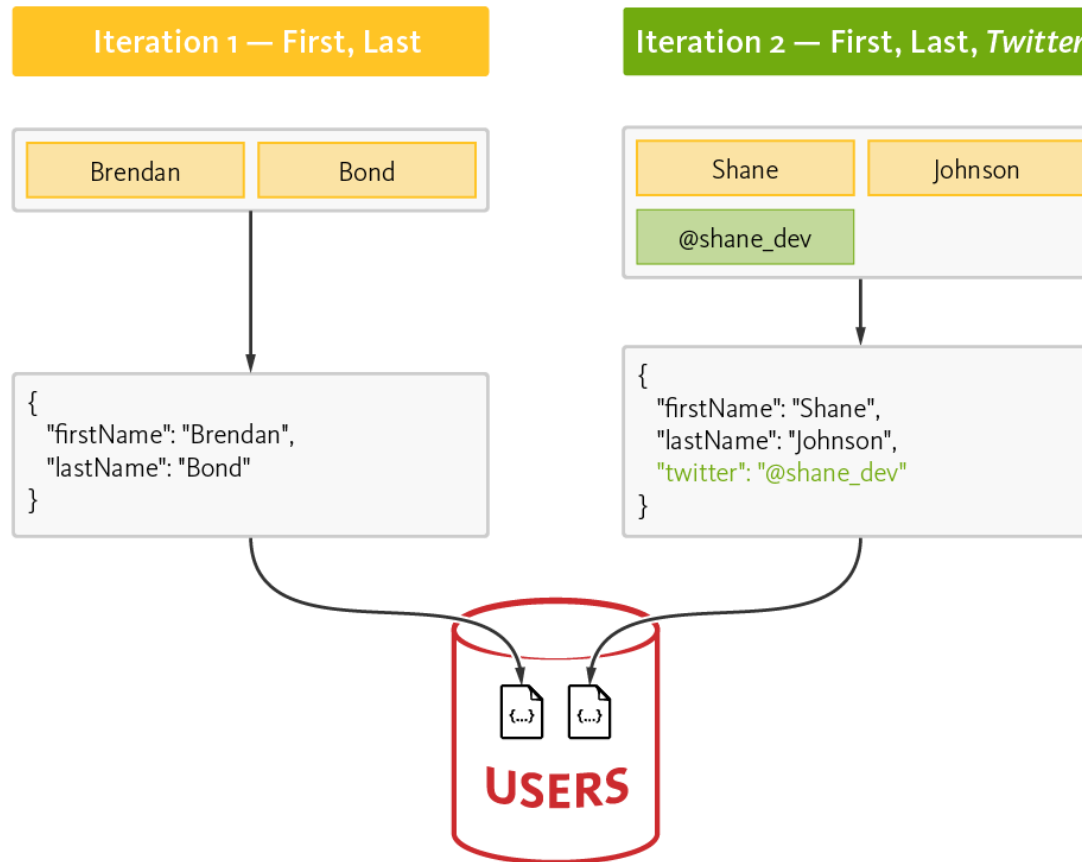
USERS		
0	Brendan	Bond
1	Shane	Johnson

This can not be stored

# NoSQL - Flexibility



In NoSQL DB, we can easily add new columns.

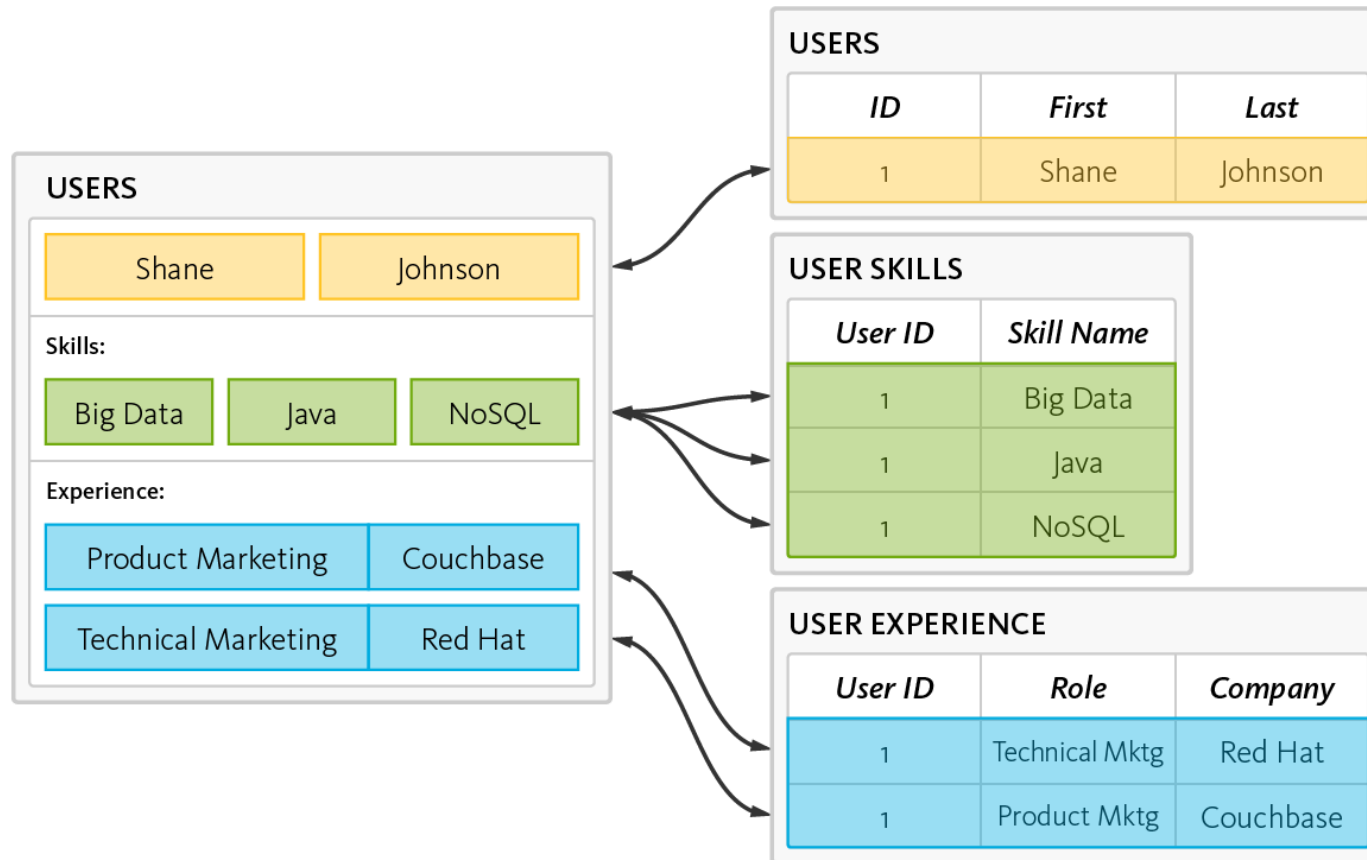


Ref: <https://www.couchbase.com/resources/why-nosql>

November 9, 2024

SE 70681/US 70678 Software Architectures

# NoSQL - Simplicity



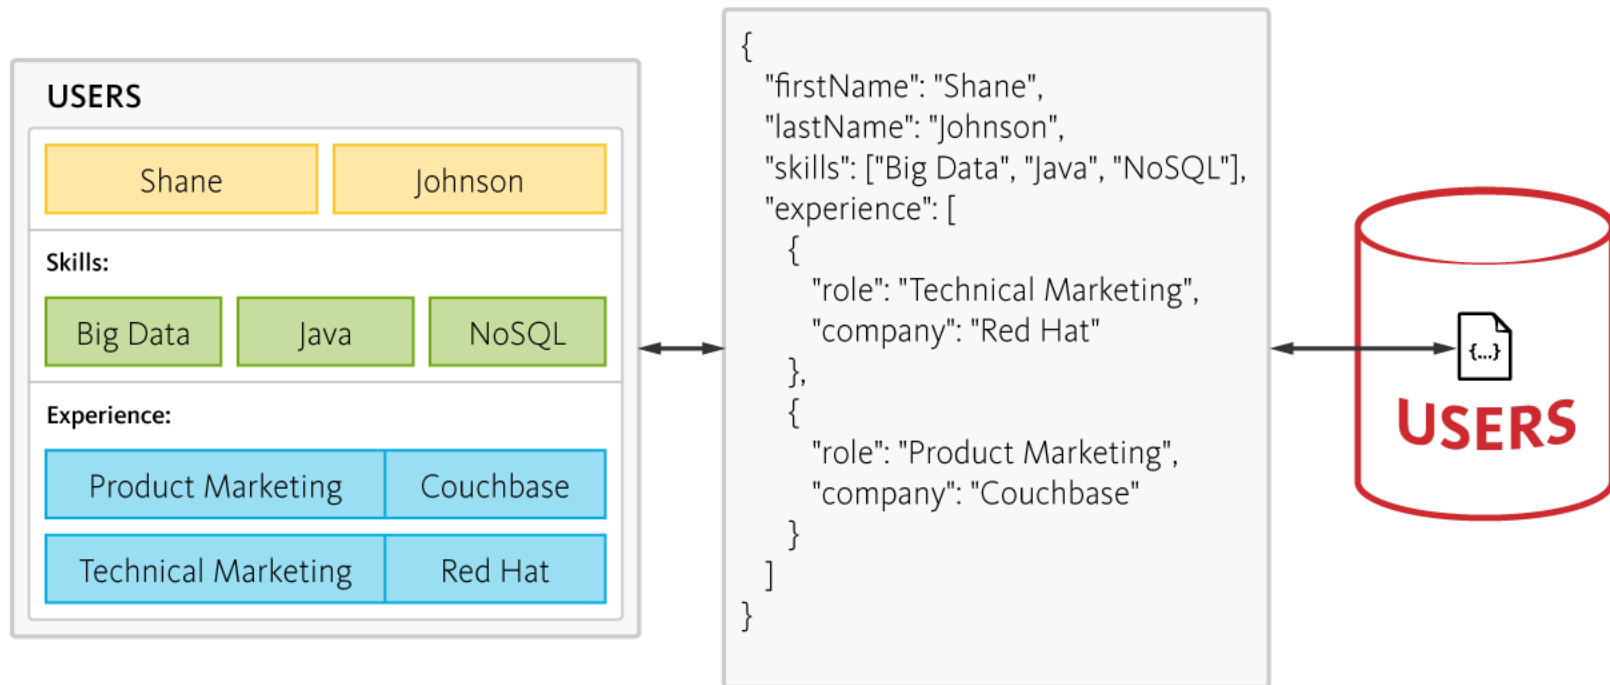
NoSQL

November 9, 2024

SE ZG651/ SS ZG653 Software  
Architectures

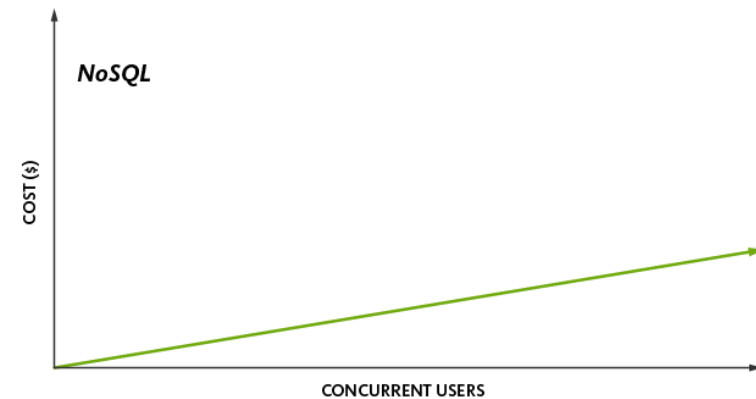

SQL

# NoSQL - Simplicity





# NoSQL – Cost effective



Cost: Memory, CPU, storage

# NoSQL Database



Types of NoSQL databases:

- Document
- Key Value
- Column stores
- Graph

# Document NoSQL database



- The database stores and retrieves documents, which can be XML, JSON, BSON, and so on.
- Some of the popular document databases we have seen are [MongoDB](#), [CouchDB](#), [Terrastore](#), [OrientDB](#), [RavenDB](#), and of course the well-known and often reviled Lotus Notes that uses document storage.

```
<Key=CustomerID>

{
  "customerid": "fc986e48ca6" ← Key
  "customer":
  {
    "firstname": "Pramod",
    "lastname": "Sadhalage",
    "company": "ThoughtWorks",
    "likes": [ "Biking", "Photography" ]
  }
  "billingaddress":
  { "state": "AK",
    "city": "DILLINGHAM",
    "type": "R"
  }
}
```

Example of one record

---

RDBMS- single node for read and write

- Consistent
- Not partition tolerant.
- Compromise availability during write.

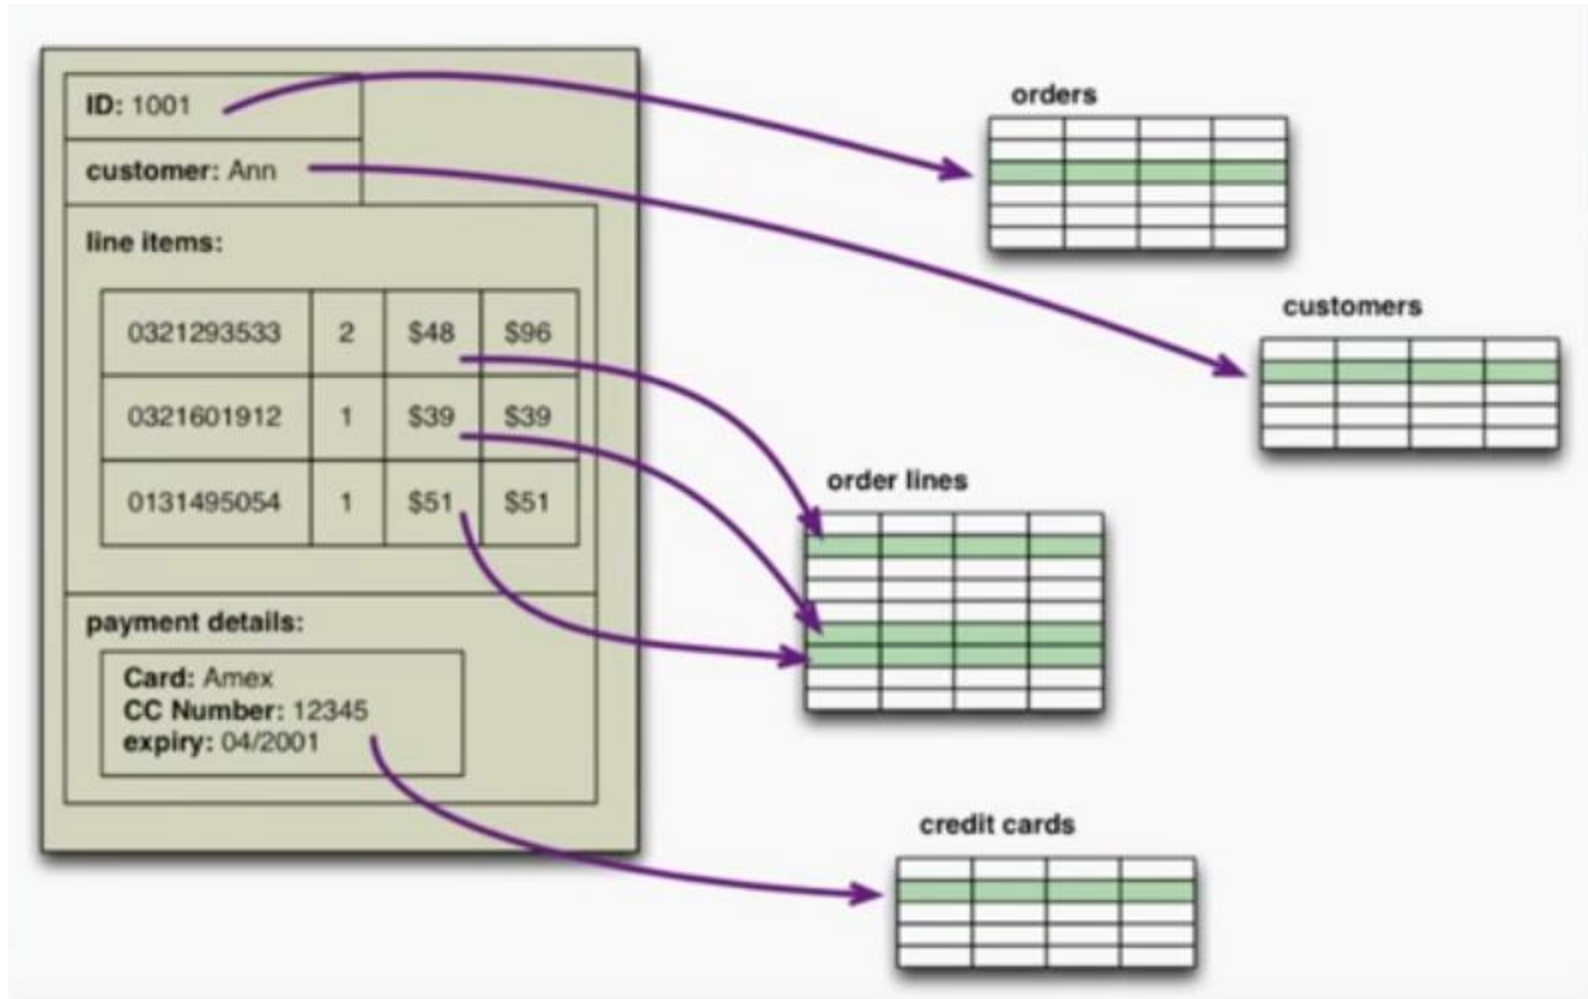
Two Machines:

- High Availability
- Partition Tolerant
- No Consistency.

One machine with backup – HDFS/RDBMS with mirror

- Consistent, Partition tolerant but not highly available.

# Document DB vs Relational DB



---

Good for

- Ecommerce platform
- Content management systems

# Features of Mongo DB

---

- Indexing
- Ad hoc search
- Replication
- Partitioning / Sharding
  - Ex. Partition data by Product or Geography

# Key-Value database

- Data model is Key value pair
- **Uses hashing for fast access**
- The DB does not care what is contained in the value
- Example scenarios: Phone directory, Stock trading
- Some of the popular key-value databases are [Riak](#), [Redis](#) (often referred to as Data Structure server), [Memcached](#) and its flavors, [Berkeley DB](#), [upscaledb](#) (especially suited for embedded use), Amazon DynamoDB (not open-source), Project Voldemort and [Couchbase](#).

## Phone Directory

Key	Value
Bob	(123) 456-7890
Jane	(234) 567-8901
Tara	(345) 678-9012
Tiara	(456) 789-0123



# Example of key value database



## Stock Trading

This example uses a list as the value.

The list contains the stock ticker, whether its a “buy” or “sell” order, the number of shares, and the price.

Key	Value
123456789	APPL, Buy, 100, 84.47
234567890	CERN, Sell, 50, 52.78
345678901	JAZZ, Buy, 235, 145.06
456789012	AVGO, Buy, 300, 124.50

More examples: User profiles, Blog comments

November 9, 2024

SE ZG651/ SS ZG653 Software  
Architectures

17

# Uses of Redis



- Session cache, with persistence

# Column oriented database



This is useful for data analysis scenarios

Example: Calculate the average usage of electricity in 2018 in East Bangalore region

Traditional way: Read all the billing records of 2018

<u>Cust id,</u>	<u>Name,</u>	<u>Addr,</u>	<u>Region,</u>	<u>Month,</u>	<u>Year,</u>	<u>Usage,</u>	<u>Amt,</u>
Record 1: 001,	John Mancha,	Addr 1,	East,	Jan,	2018,	100,	600
Record 2: 002,	Vivek Kulkarni,	Addr 2,	East,	Jan,	2018,	90,	540
Record 3: 003,	Shanti Sharma,	Addr 3,	West,	Jan,	2018,	110,	660
....							
....							

# Column oriented database



Instead if we store the data as follows:

Record 1: 001, John Mancha, Addr 1, East // Customer details

Record 2: 002, Vivek Kulkarni, Addr 2, East

Record A: Jan, 2018, 100, 90, 110, ... // Usage – in customer order

Record B: Feb, 2018, 110, 92, 115, ...

Only one record 'Record A' is needed to calculate average usage in Jan 2018

Good for summarization

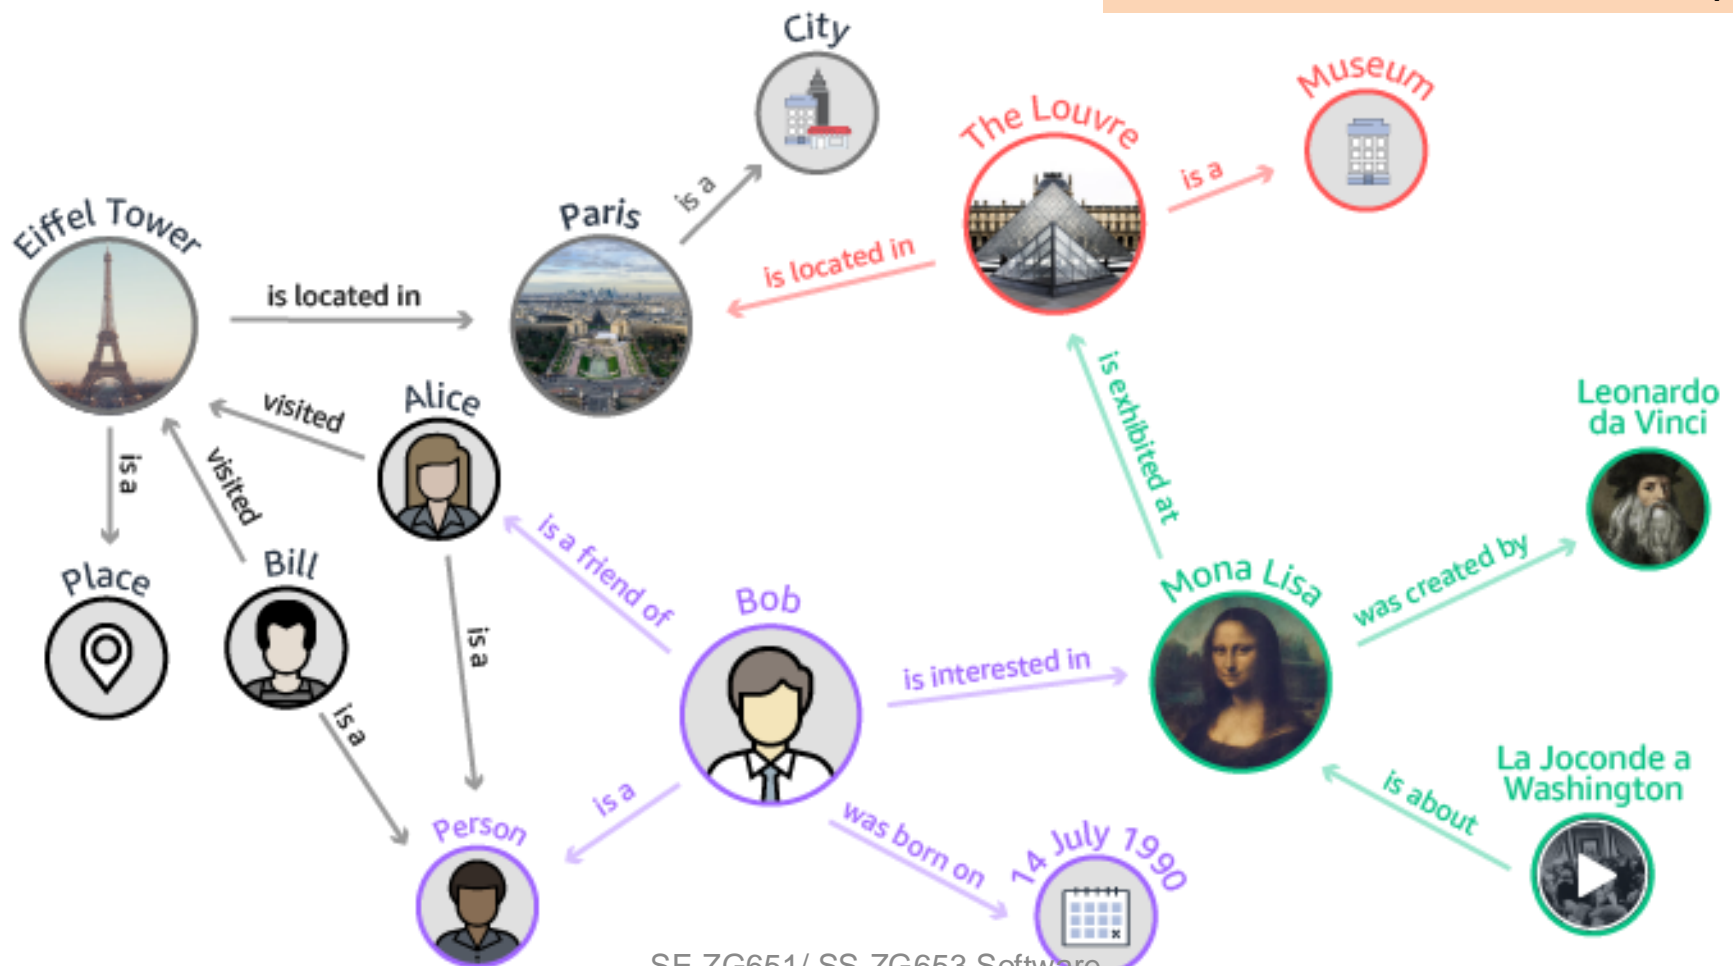
# Graph database: Example: Knowledge graph

innovate

achieve

lead

Entities and their relationships



November 9, 2024

SE ZG651/ SS ZG653 Software  
Architectures

21

# Graph database



Entities and relationships have properties (attributes)

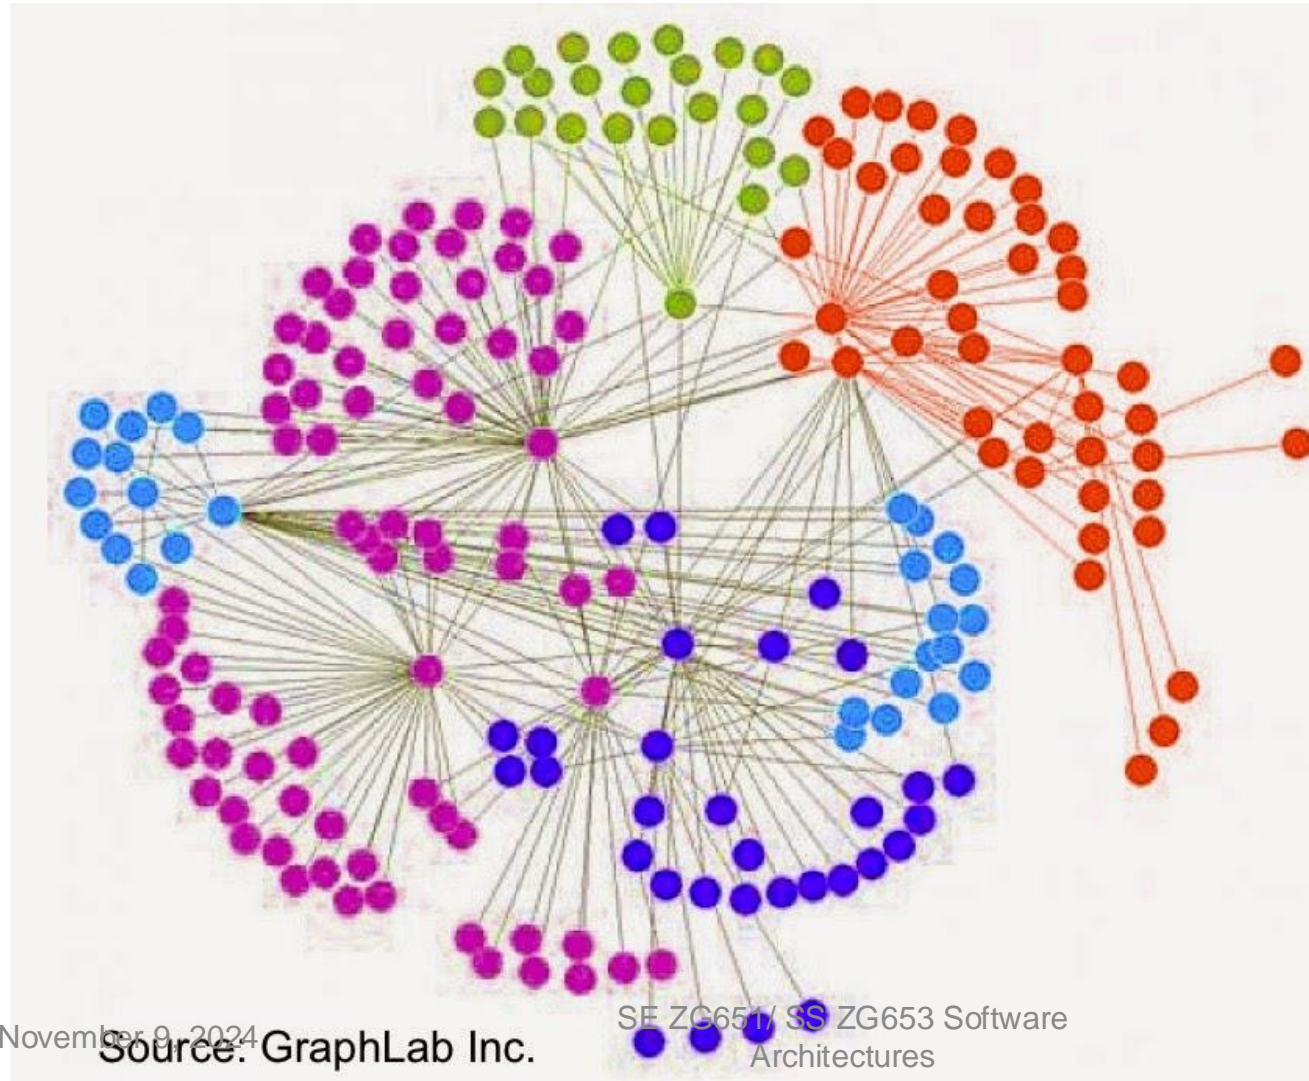
Ex.

- Eifel tower properties can be height, date of construction
- Visited relationship can have properties such as date of visit

## Uses

- Store a large amount of inter-related information and to search for an entity along with its relationships
- Fraud detection
- Social networks

# Can be use to detect hidden patterns



Easy to understand clusters

November 9, 2024

Source: GraphLab Inc.

SE ZG651/ SS ZG653 Software Architectures

23

# Graph database



- Graph databases allow you to store entities and relationships between these entities.
- Entities are also known as nodes, which have properties. Think of a node as an instance of an object in the application.
- Relations are known as edges that can have properties. Edges have directional significance; nodes are organized by relationships which allow you to find interesting patterns between the nodes.
- There are many graph databases available, such as [Neo4J](#), [Infinite Graph](#), [OrientDB](#), or [FlockDB](#)



# In-Memory databases

- An **in-memory database (IMDB; also main memory database system or MMDB or memory resident database)** is a [database management system](#) that primarily relies on [main memory](#)
- Applications where response time is critical, such as those running telecommunications network equipment and [mobile advertising](#) networks, often use main-memory databases
- With the introduction of [non-volatile random access memory](#) technology (Flash memory), in-memory databases will be able to run at full speed and maintain data in the event of power failure
- Popular In-memory databases are **SAP's HANA**, IBM DB2 BLU, Oracle
- **These databases support OLTP and OLAP (Online Analytical Processing)**

# Acknowledgement

---



## Sources

### Hadoop

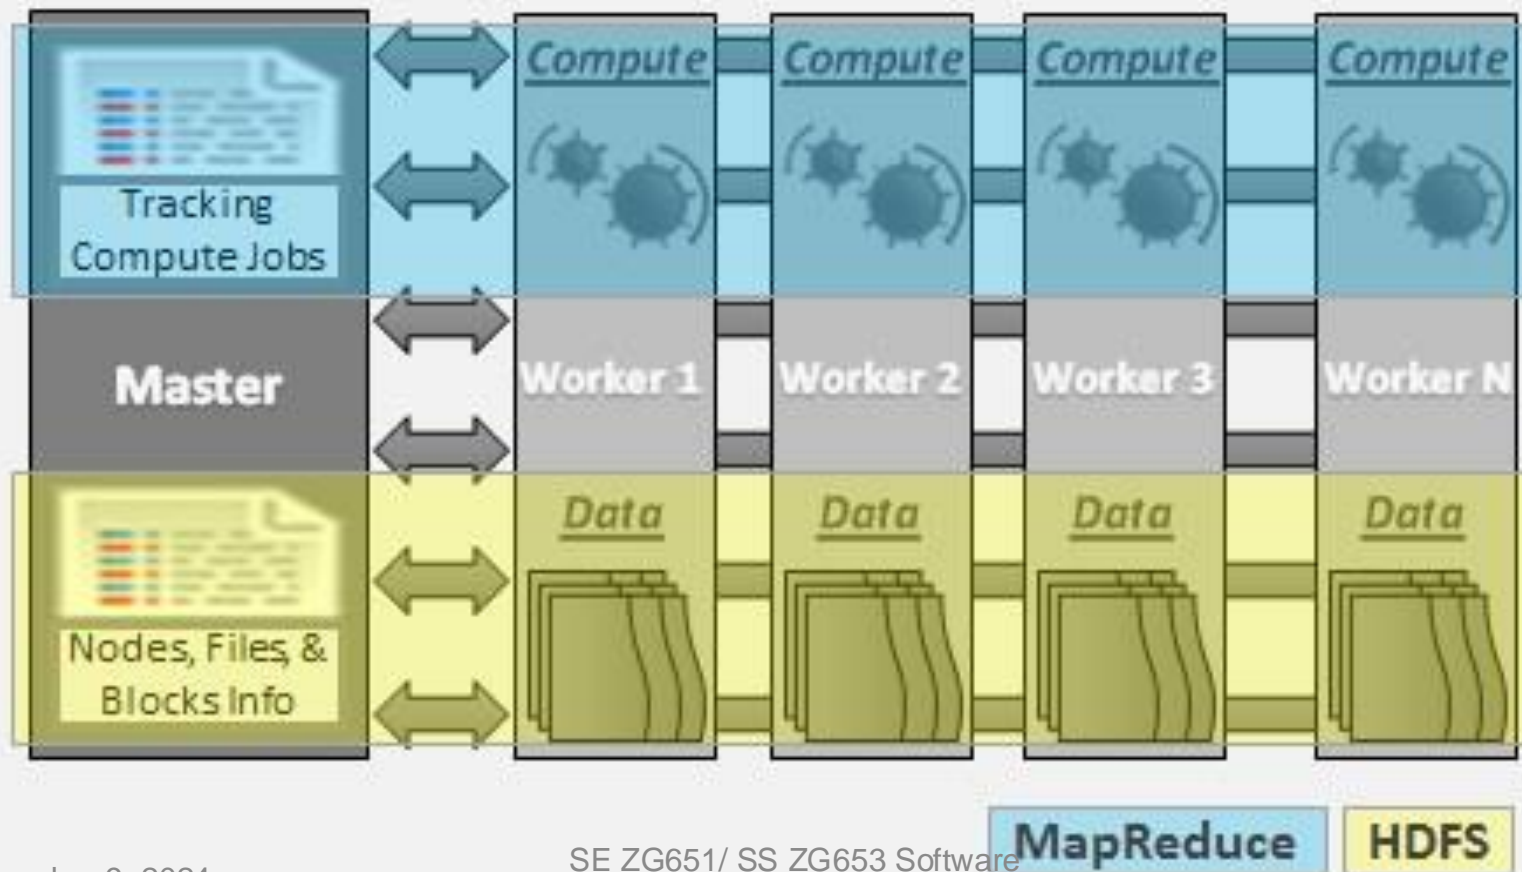
<https://www.mssqltips.com/sqlserverauthor/77/dattatreysindol/>

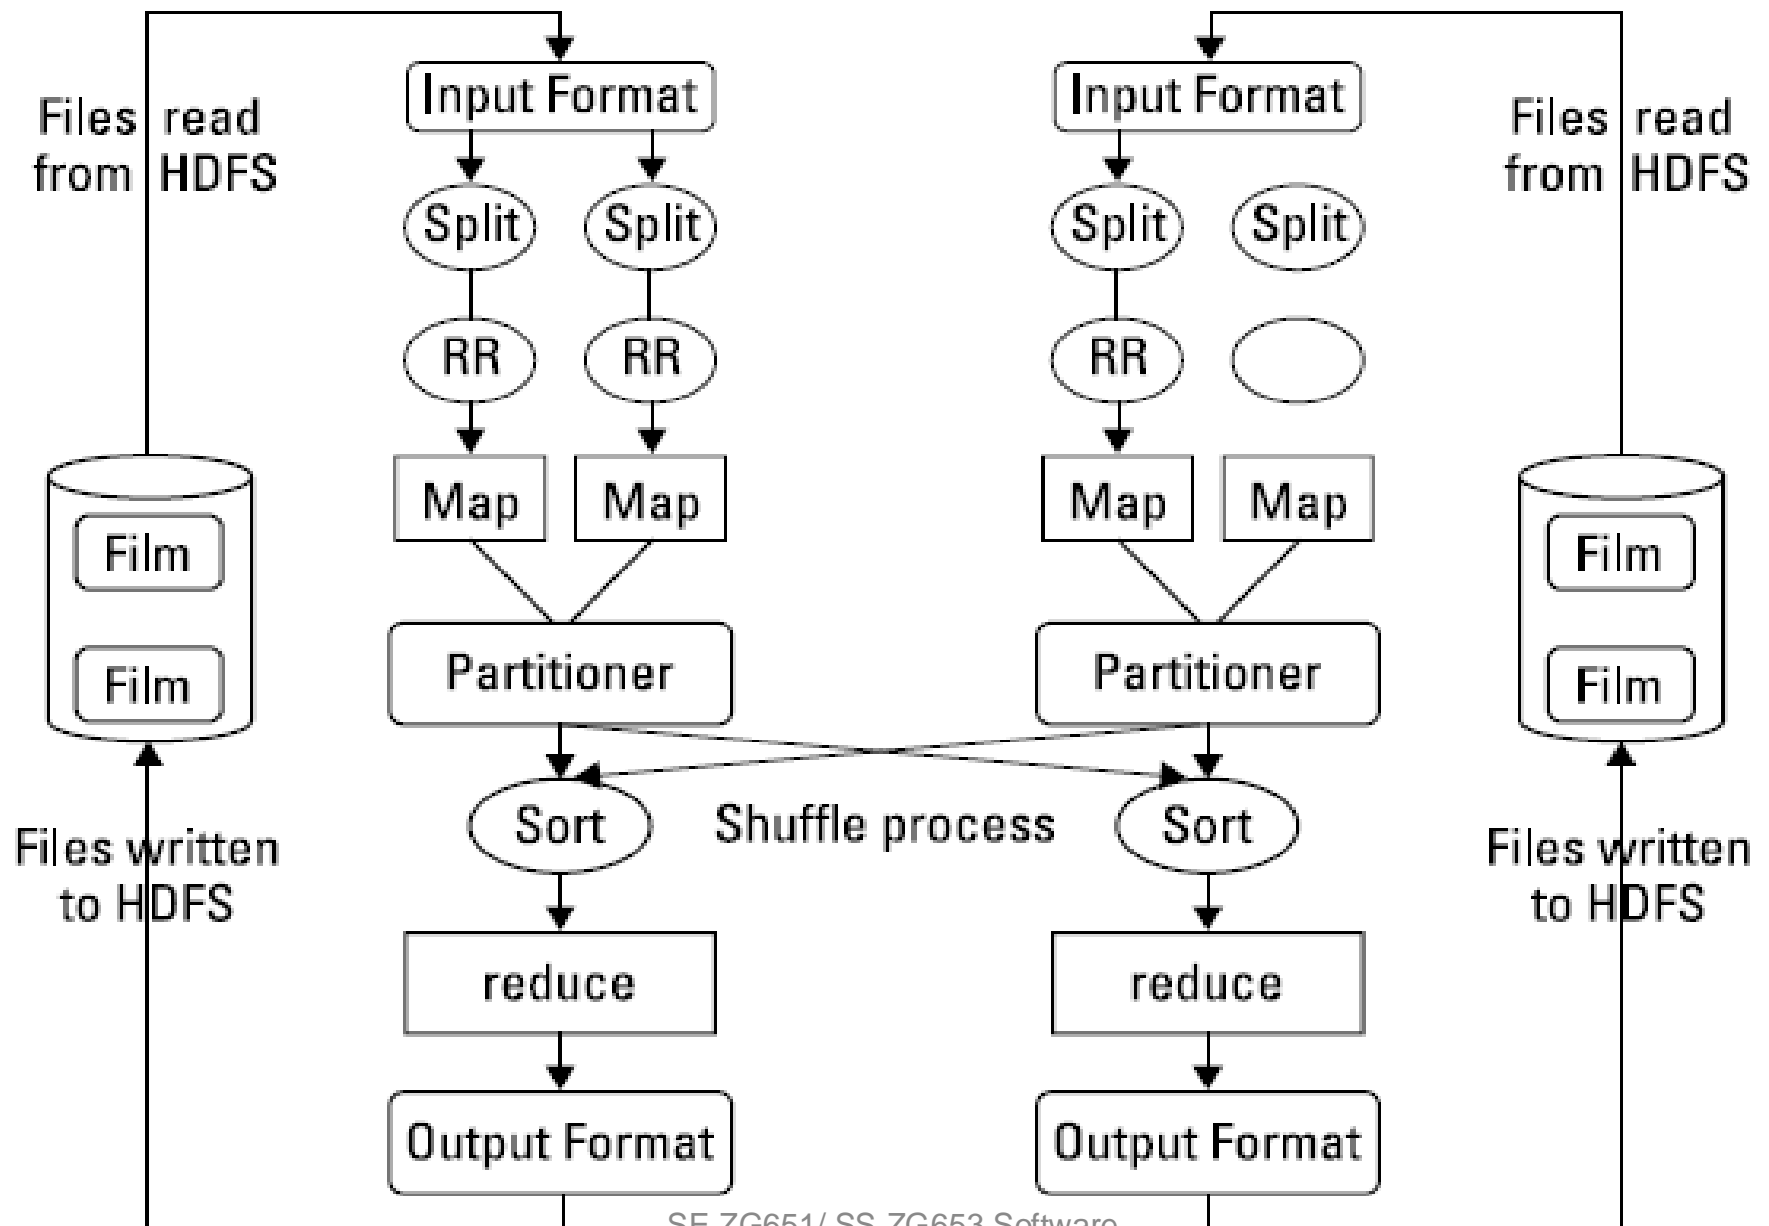
### NoSQL Database

<https://www.thoughtworks.com/insights/blog/nosql-databases-overview>

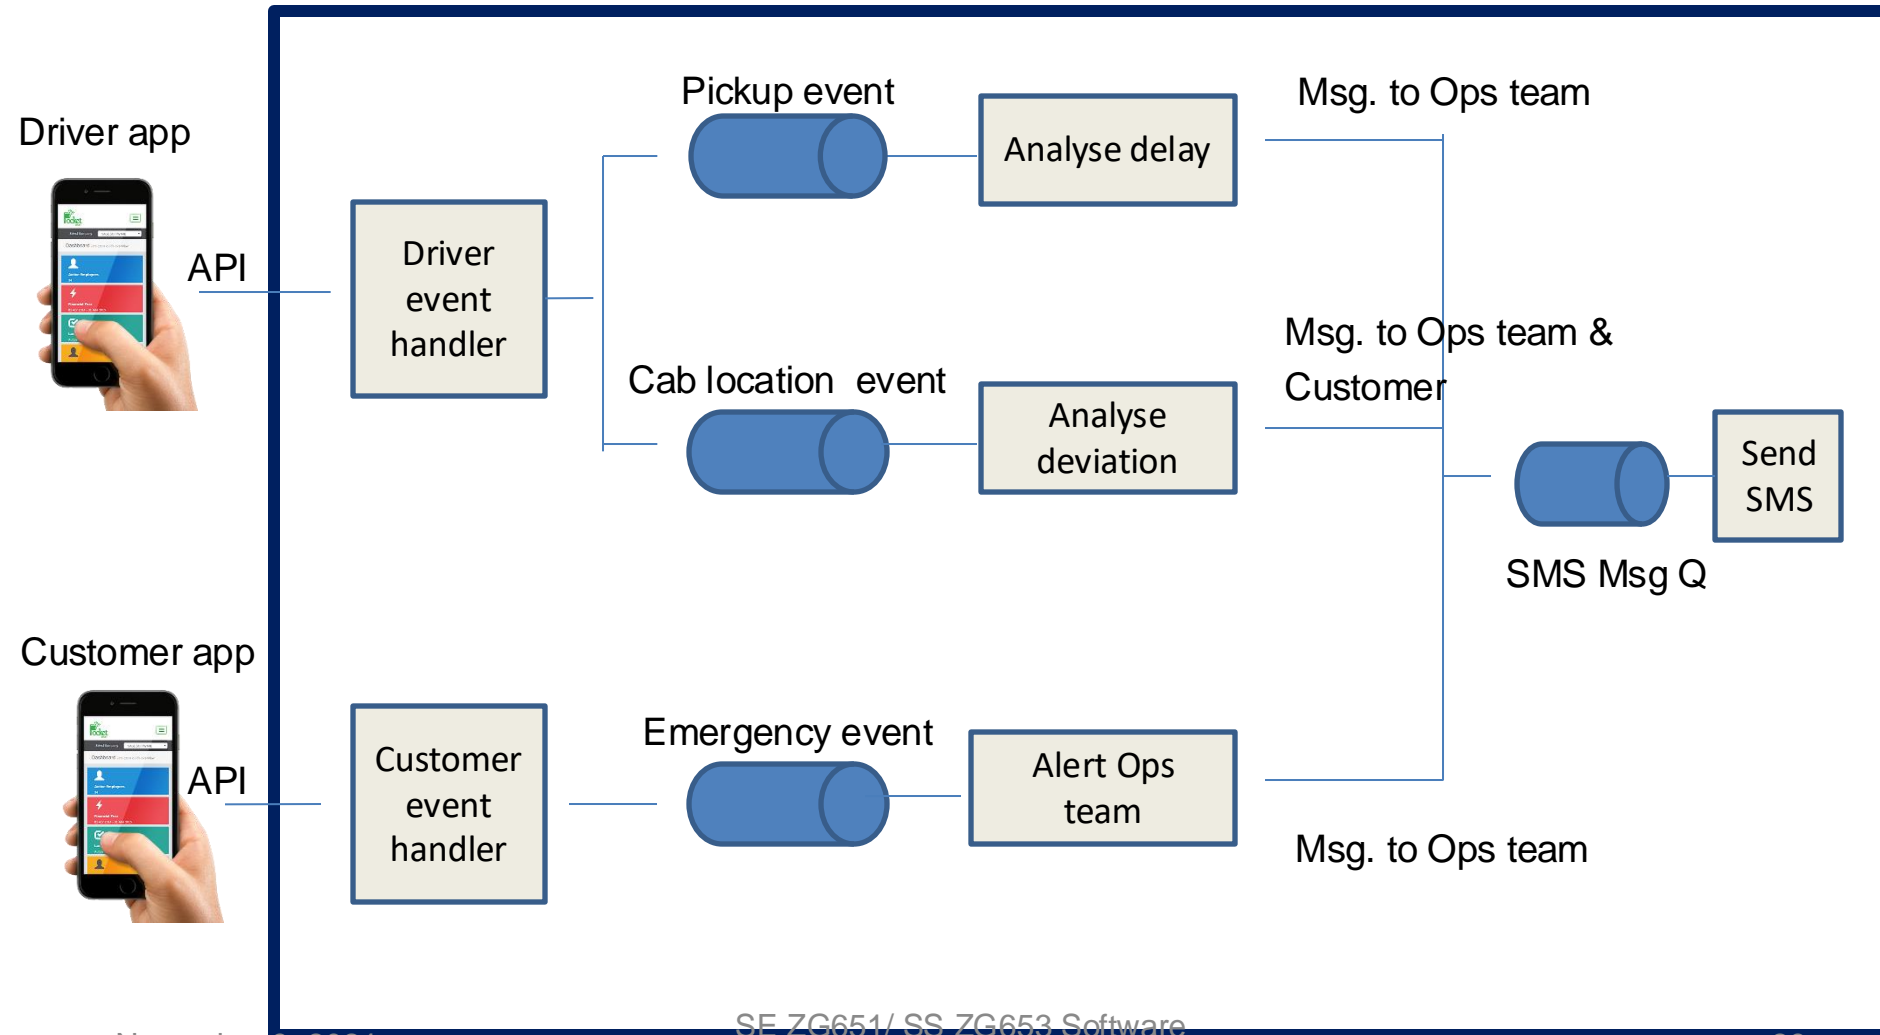
# Typical Architecture of Hadoop

Across Same or Different Rack, Data Center, or Region





# Example of Stream processing in a cab hailing company like Ola / Uber



# Analytics



Data Visualization

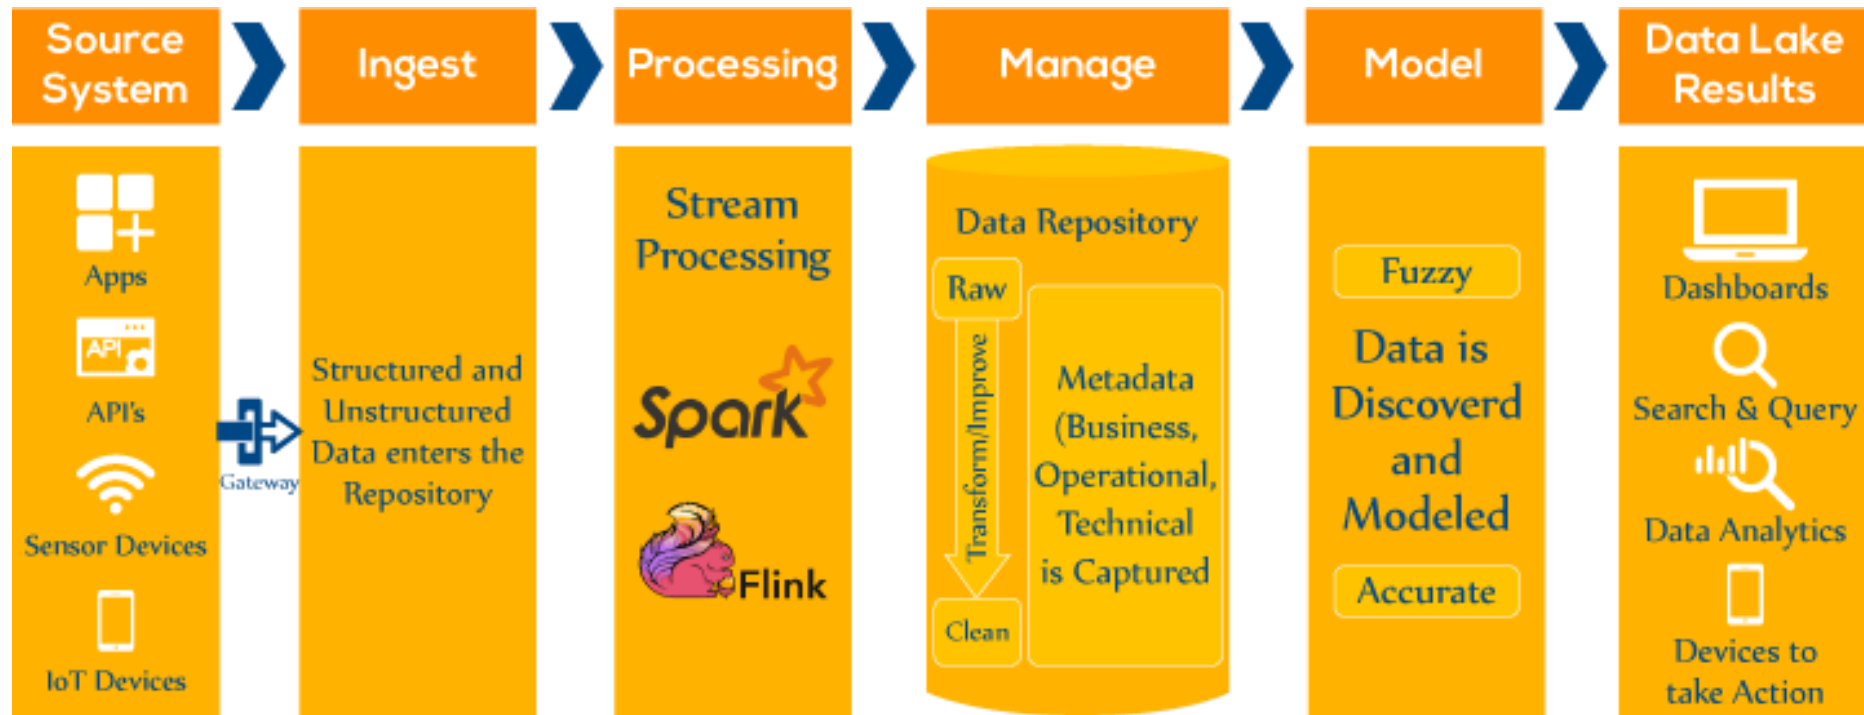
Multi-dimensional data

Data mining

Examples...

Tools ....

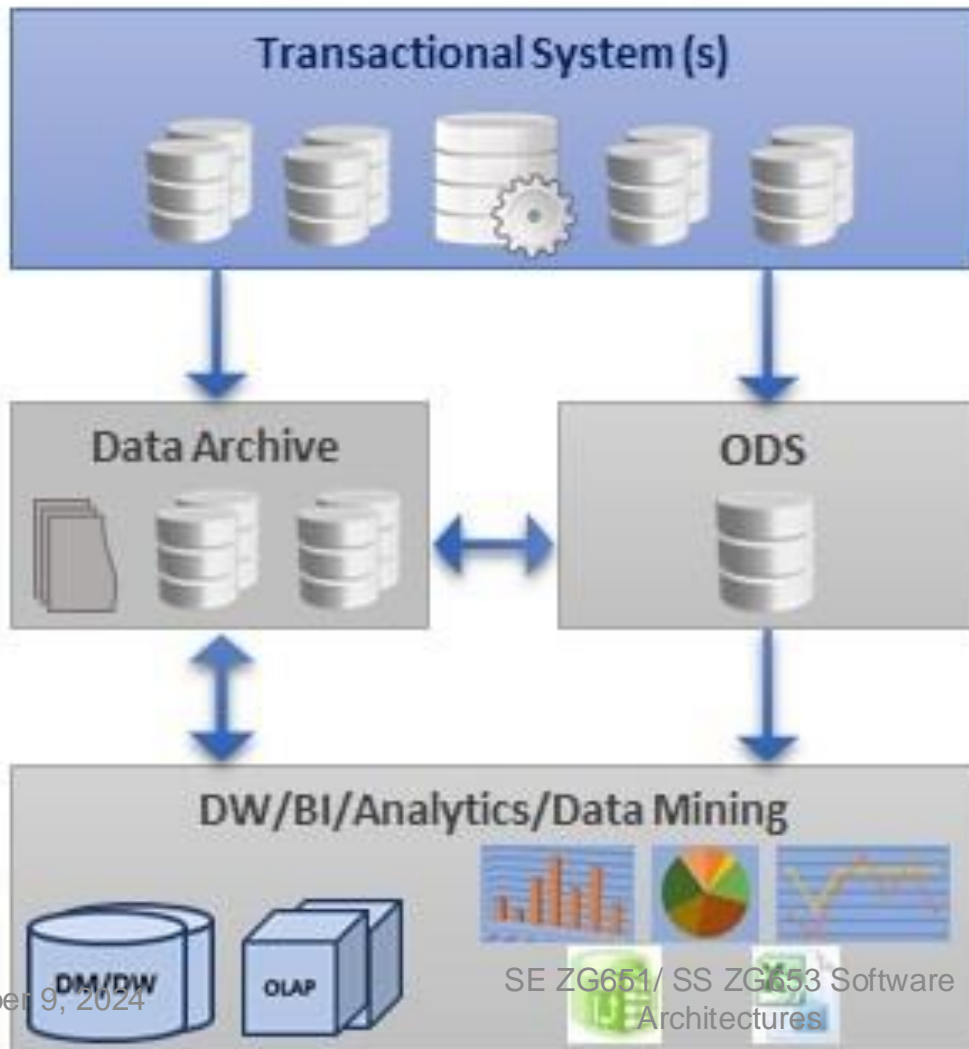
# Real-Time Streaming and Data Analytics For IoT



# Traditional Data Processing



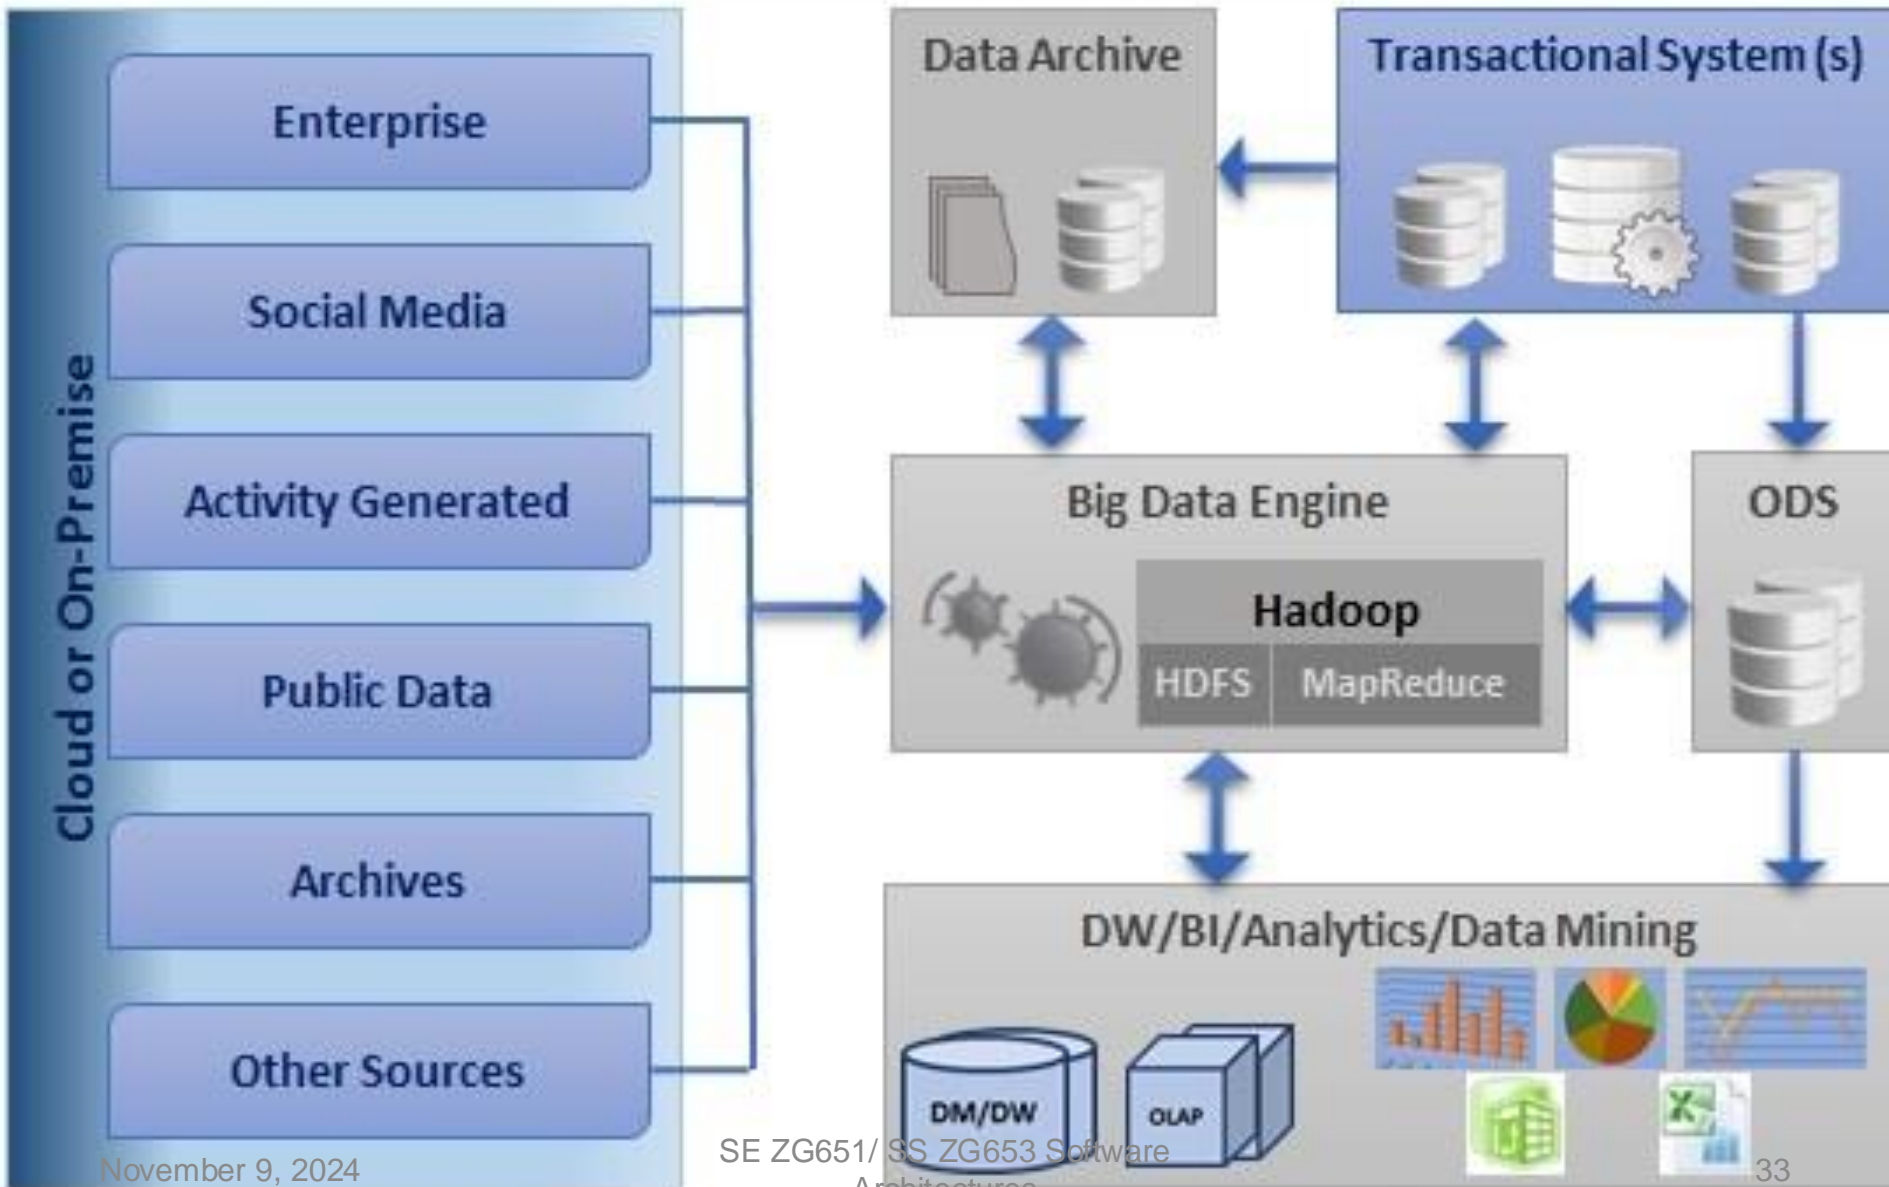
## Traditional Data Processing & Management



ODS: Operational  
Data Store



## Modern (Next Generation) Data Processing & Management



November 9, 2024

SE ZG651/ SS ZG653 Software  
Architectures

33

# Use cases of Stream Processing



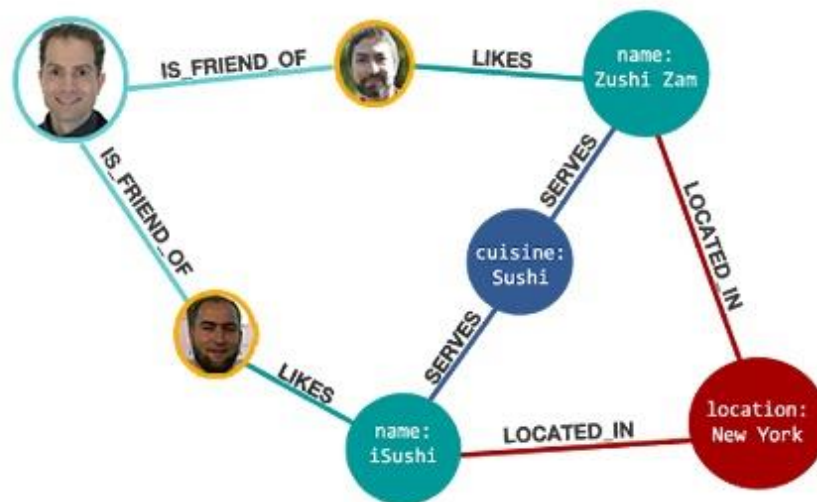
Following are some of the use cases.

- Algorithmic Trading, **Stock Market Surveillance**,
- Smart Patient Care
- **Monitoring a production line**
- Supply chain optimizations
- Intrusion, Surveillance and Fraud Detection ( e.g. [Uber](#))
- Most Smart Device Applications : Smart Car, Smart Home ..
- **Smart Grid**—(e.g. load prediction and outlier plug detection see [Smart grids, 4 Billion events, throughout in range of 100Ks](#))
- Traffic Monitoring, Geo fencing, Vehicle and Wildlife tracking—e.g. [TFL London Transport Management System](#)
- Sport analytics—Augment Sports with realtime analytics (e.g. this is a work we did with a real football game (e.g. [Overlaying realtime analytics on Football Broadcasts](#))
- Context-aware promotions and advertising
- Computer system and network monitoring
- Predictive Maintenance, (e.g. [Machine Learning Techniques for Predictive Maintenance](#))
- Geospatial data processing

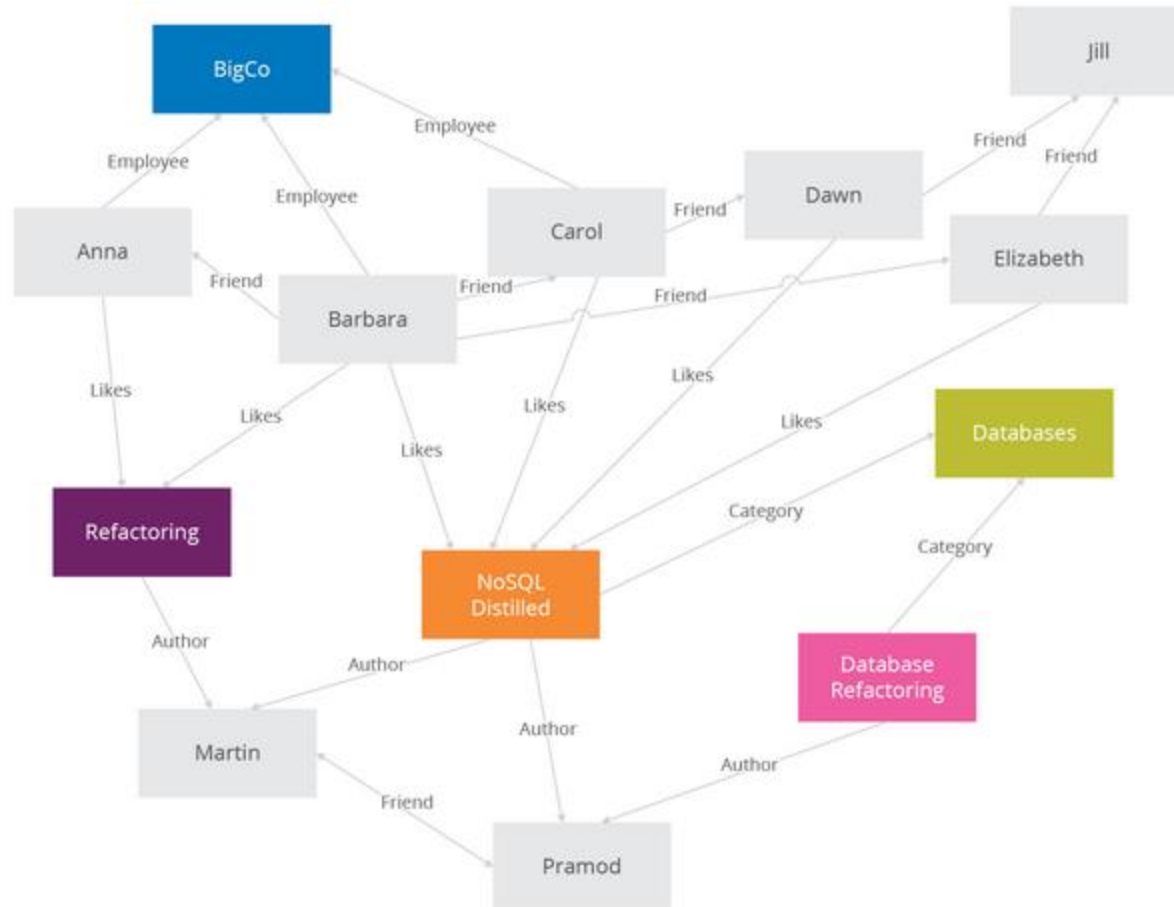
November 9, 2024

SE ZG651/ SS ZG653 Software  
Architectures


34




# Graph database



# Difference between Hadoop & Spark...



Interactive operation using Hadoop



Interactive operation using Spark

Source:  
[https://www.tutorialspoint.com/apache\\_spark/apache\\_spark\\_rdd.htm](https://www.tutorialspoint.com/apache_spark/apache_spark_rdd.htm)

# Real-time analytics



Detecting bank fraud requires real-time analytics as events happen

Such situations demand processing of each event as they happen rather processing a batch of data on disk

This led to tools such as Spark Streams and Storm which support in-memory processing, than disk based processing

# Storm for real time computing



Apache Storm is a free and open source **distributed real time computation system**.



Storm makes it easy to reliably process unbounded streams of data, doing for realtime processing what Hadoop did for batch processing. Storm is simple, can be used with any programming language, and is a lot of fun to use!

Storm has many **use cases: real time analytics, online machine learning, continuous computation**, distributed RPC, ETL, and more. Storm is fast: a benchmark clocked it at over a million tuples processed per second per node. It is scalable, fault-tolerant, guarantees your data will be processed, and is easy to set up and operate.

Storm **integrates with the queueing and database technologies** you already use. A Storm topology consumes streams of data and processes those streams in arbitrarily complex ways, repartitioning the streams between each stage of the computation however needed. Read more in the tutorial.

Source: <http://storm.apache.org/>

# Storm topology



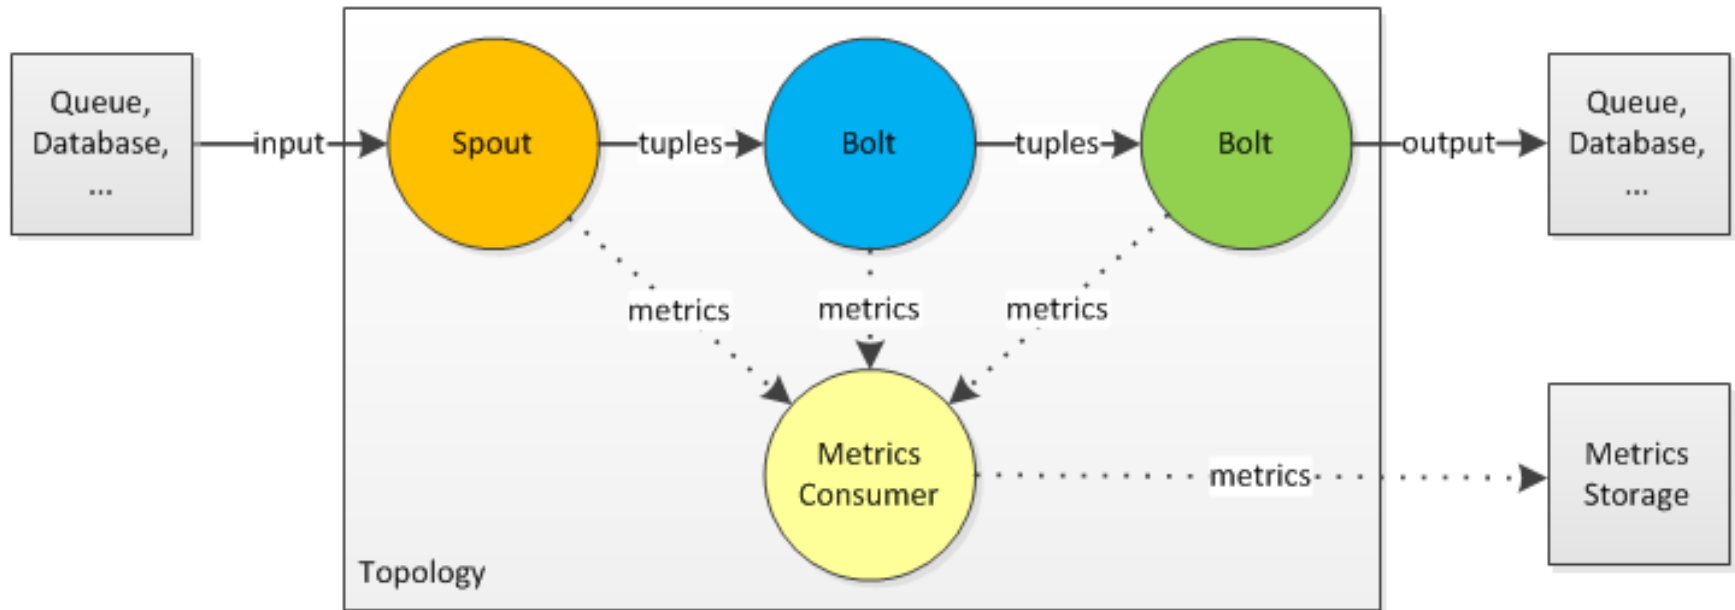
Twitter uses Storm for real-time analytics, personalization, search, revenue optimization

Groupon uses Storm for Real-time data integration systems

Yahoo! Uses Storm for processing user events, content feeds, and application logs.



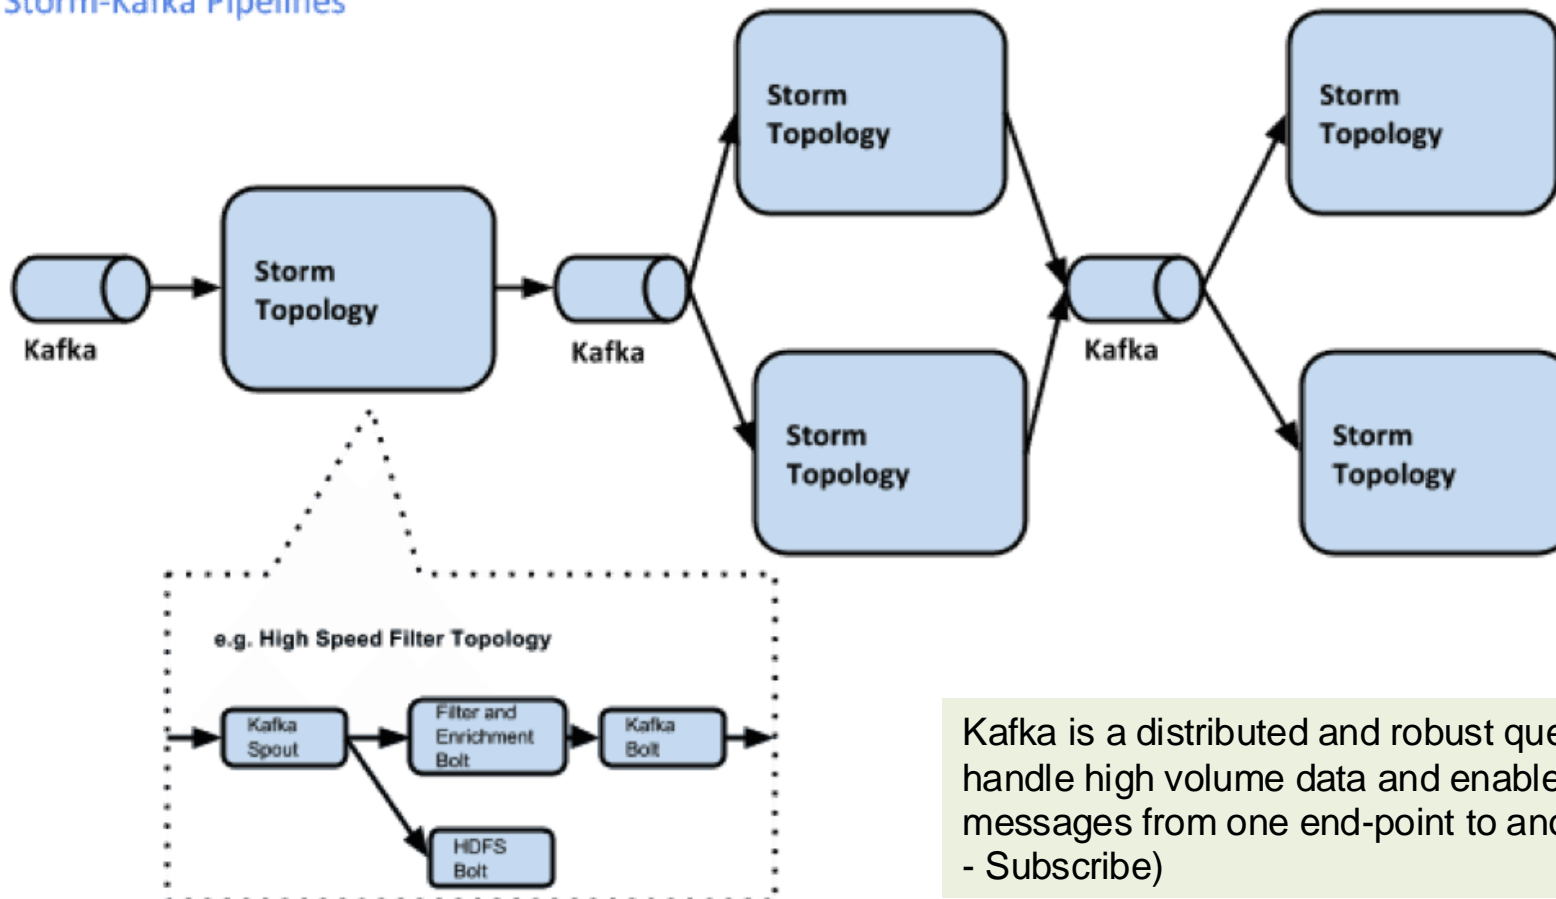
# Architecture of Storm system



# Combining Kafka and Storm for real time computing



## Storm-Kafka Pipelines



Kafka is a distributed and robust queue that can handle high volume data and enables you to pass messages from one end-point to another (Publish - Subscribe)

# Storm & Kafka



The common flow of these tools (as I know it) goes as follows:

real-time-system --> Kafka --> Storm --> NoSql --> BI(optional)

# MapReduce – Word Count Example Flow

