

Applied Machine Learning

Assignment 2

Submitted By
Mallidi Akhil Reddy
2024TM93056

Problem Definition:

The goal of this assignment is to perform a classification analysis on two real-world datasets obtained from the UCI Machine Learning Repository:

Banknote Authentication Dataset: The goal here is to classify whether a banknote is authentic or not.

Haberman's Survival Dataset: The goal here is to predict patient survival status based on surgery data.

The following machine learning algorithms are employed for both tasks:

- 1) Naive Bayes
- 2) Logistic Regression
- 3) Support Vector Machine (SVM)
- 4) Random Forest

Mathematical Details:

Naive Bayes: Naive Bayes is probabilistic classifier algorithm based on Bayes' Theorem. It assumes independence among features.

$$P(C|X) = \frac{P(X|C)P(C)}{P(X)}$$

Logistic Regression: It's a linear model algorithm for binary classification. In this Model parameters are estimated using maximum likelihood estimation.

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Support Vector Machine (SVM): This algorithm Maximizes the margin between classes.

$$\begin{array}{ll} \min_{\mathbf{w}, b} & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to:} & y^i (\mathbf{w}^T \mathbf{x}^i + b) \geq 1 \quad \forall i \end{array}$$

Dataset Description:

Banknote Authentication Dataset:

- Number of features: 4 (variance, skewness, kurtosis, entropy)
- Number of instances: 1372
- Number of classes: 2 (authentic, counterfeit)
- Preprocessing steps:
 - 1) Scaling the features using StandardScaler to have zero mean and unit variance.
This can help improve the performance of algorithms like Logistic Regression and SVM.
 - 2) No missing values were found in the dataset.

Haberman's Survival:

- Number of features: 3 (age, year of operation, number of positive axillary nodes)
- Number of instances: 306
- Number of classes: 2 (survived 5 years or more, died within 5 years)
- Preprocessing steps:
 - 1) Scaling the features using StandardScaler, as it is important because the features have different ranges, and scaling can prevent features with larger ranges.
 - 2) No missing values were found.

Experiments:

Experimental Setup and Train-Test Split:

For each dataset, I used an 80/20 train-test split. This means that 80% of the data was used for training the models, and 20% was held out for evaluating their performance.

Hyperparameter Tuning:

Naive Bayes: Naive Bayes generally doesn't have hyperparameters that require extensive tuning.

Logistic Regression:

- Hyperparameter: Regularization strength (C)
- Tuning method: Grid search
- Parameter range: C = [0.001, 0.01, 0.1, 1, 10, 100]
- Regularization type: l2

SVM:

- Hyperparameters: Kernel type, C (regularization), gamma (for RBF kernel)
- Tuning method: Grid search
- Parameter ranges:
 - Kernel: ['linear', 'rbf']
 - C: [0.1, 1, 10, 100]
 - gamma (for RBF): [0.001, 0.01, 0.1, 1]

Results on Test Data:

Metrics are rounded to 4 decimals

Banknote Authentication Dataset:

Model	Accuracy	Macro F1-Score	Macro Precision	Macro Recall
Naive Bayes	0.8582	0.856	0.8571	0.8551
Logistic Regression	0.9855	0.9853	0.9841	0.9869
Support Vector Machine (SVM)	1.0	1.0	1.0	1.0

Haberman Survival Dataset:

Model	Accuracy	Macro F1-Score	Macro Precision	Macro Recall
Naive Bayes	0.7581	0.57	0.686	0.572
Logistic Regression	0.7581	0.5338	0.7147	0.5516
Support Vector Machine (SVM)	0.7258	0.5127	0.5772	0.5299

Conclusion:

The classification of the Banknote Authentication dataset yielded remarkably high performance across all evaluated models. Support Vector Machine (SVM) achieved a perfect score with 100% accuracy, F1-Score, precision, and recall on the test data. Logistic Regression also demonstrated strong performance with an accuracy of 0.9855, a macro F1-Score of 0.9853, a macro precision of 0.9841, and a macro recall of 0.9869. Naive Bayes, while slightly lower, still provided a respectable accuracy of 0.8582, a macro F1-Score of 0.856, a macro precision of 0.8571, and a macro recall of 0.8551.

The classification of the Haberman's Survival dataset proved to be a more challenging task. All three models exhibited lower performance compared to the banknote dataset. Logistic Regression and Naive Bayes achieved the same accuracy of 0.7581. However, Logistic Regression showed a slightly better macro F1-Score (0.5338) compared to Naive Bayes (0.57), although Naive Bayes had a higher macro precision (0.686 vs 0.7147 for Logistic Regression) but a lower macro recall (0.572 vs 0.5516 for Logistic Regression). SVM performed slightly worse on this dataset with an accuracy of 0.7258 and a macro F1-Score of 0.5127.

Future Work:

Incorporate feature selection or transformation like PCA to explore latent patterns.

Employing stratified k-fold cross-validation to ensure stable performance estimates across imbalanced data splits.

Code

```

# importing all required libraries

import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
recall_score

# Load the datasets

# Banknote Authentication Dataset

banknote_data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-
learning-databases/00267/data_banknote_authentication.txt',
header=None, names=['variance', 'skewness', 'curtosis', 'entropy',
'class'])
X_banknote = banknote_data.drop('class', axis=1)
y_banknote = banknote_data['class']

print("\n Banknote Dataset Description:")
print(f" Number of features: {X_banknote.shape[1]}")
print(f" Number of instances: {X_banknote.shape[0]}")
print(f" Number of classes: {len(y_banknote.unique())}")
print(" Features:", list(X_banknote.columns))

Banknote Dataset Description:
Number of features: 4
Number of instances: 1372
Number of classes: 2
Features: ['variance', 'skewness', 'curtosis', 'entropy']

# Haberman Dataset

haberman_data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-
learning-databases/haberman/haberman.data', header=None, names=['age',
'year', 'nodes', 'survival'])
X_haberman = haberman_data.drop('survival', axis=1)
y_haberman = haberman_data['survival'].map({1: 1, 2: 0})

# Dataset Description
print("\n Haberman Dataset Description:")
print(f" Number of features: {X_haberman.shape[1]}")
print(f" Number of instances: {X_haberman.shape[0]}")
print(f" Number of classes: {len(y_haberman.unique())}")
print(" Features:", list(X_haberman.columns))

```


Haberman Dataset Description:

Number of features: 3

Number of instances: 306

Number of classes: 2

Features: ['age', 'year', 'nodes']

Preprocessing datasets

```
scaler_banknote_data = StandardScaler()
X_banknote_scaled_data =
scaler_banknote_data.fit_transform(X_banknote)
```

```
scaler_haberman_data = StandardScaler()
X_haberman_scaled_data =
scaler_haberman_data.fit_transform(X_haberman)
```

Train-test split with 80% train data and 20% test data

```
X_banknote_train, X_banknote_test, y_banknote_train, y_banknote_test =
train_test_split(X_banknote_scaled_data, y_banknote, test_size=0.2,
random_state=42, stratify=y_banknote)
X_haberman_train, X_haberman_test, y_haberman_train, y_haberman_test =
train_test_split(X_haberman_scaled_data, y_haberman, test_size=0.2,
random_state=42, stratify=y_haberman)
```

Model Training and Evaluation

```
results = {}
```

```
results['banknote'] = {}
```

Banknote Authentication Models

1. Naive Bayes

```
print("Banknote Authentication Dataset \n")
print("Training Naive Bayes \n")
```

```
nb_banknote = GaussianNB()
nb_banknote.fit(X_banknote_train, y_banknote_train)
y_banknote_pred_nb = nb_banknote.predict(X_banknote_test)
results['banknote']['Naive Bayes'] = {
    'Accuracy': round(accuracy_score(y_banknote_test,
y_banknote_pred_nb), 4),
    'Macro F1-Score': round(f1_score(y_banknote_test,
y_banknote_pred_nb, average='macro'), 4),
    'Macro Precision': round(precision_score(y_banknote_test,
y_banknote_pred_nb, average='macro'), 4),
    'Macro Recall': round(recall_score(y_banknote_test,
y_banknote_pred_nb, average='macro'), 4)}
```

```

}

print(f"Naive Bayes Results: {results['banknote']['Naive Bayes']} \n")
print("\n-----\n")

# 2. Logistic Regression

print("Training Logistic Regression \n")

lr_banknote = LogisticRegression(random_state=42)
param_grid_lr_banknote = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
                           'solver': ['liblinear']}
grid_search_lr_banknote = GridSearchCV(lr_banknote,
                                       param_grid_lr_banknote, cv=5, scoring='f1_macro')
grid_search_lr_banknote.fit(X_banknote_train, y_banknote_train)
best_lr_banknote = grid_search_lr_banknote.best_estimator_
y_banknote_pred_lr = best_lr_banknote.predict(X_banknote_test)
results['banknote']['Logistic Regression'] = {
    'Accuracy': round(accuracy_score(y_banknote_test,
                                     y_banknote_pred_lr), 4),
    'Macro F1-Score': round(f1_score(y_banknote_test,
                                     y_banknote_pred_lr, average='macro'), 4),
    'Macro Precision': round(precision_score(y_banknote_test,
                                              y_banknote_pred_lr, average='macro'), 4),
    'Macro Recall': round(recall_score(y_banknote_test,
                                       y_banknote_pred_lr, average='macro'), 4)
}

print(f"Logistic Regression Results: {results['banknote']['Logistic Regression']} \n")
print(f"Best Logistic Regression parameters: {grid_search_lr_banknote.best_params_} \n")
print("\n-----\n")

# 3. Support Vector Machine (SVM)

print("Training SVM \n")

svm_banknote = SVC(random_state=42)
param_grid_svm_banknote = {'C': [0.1, 1, 10, 100], 'kernel':
                           ['linear', 'rbf'], 'gamma': ['scale', 'auto']}
grid_search_svm_banknote = GridSearchCV(svm_banknote,
                                       param_grid_svm_banknote, cv=5, scoring='f1_macro')
grid_search_svm_banknote.fit(X_banknote_train, y_banknote_train)
best_svm_banknote = grid_search_svm_banknote.best_estimator_
y_banknote_pred_svm = best_svm_banknote.predict(X_banknote_test)
results['banknote']['SVM'] = {
    'Accuracy': round(accuracy_score(y_banknote_test,
                                     y_banknote_pred_svm), 4),
    'Macro F1-Score': round(f1_score(y_banknote_test,

```

```

y_banknote_pred_svm, average='macro'), 4),
    'Macro Precision': round(precision_score(y_banknote_test,
y_banknote_pred_svm, average='macro'), 4),
    'Macro Recall': round(recall_score(y_banknote_test,
y_banknote_pred_svm, average='macro'), 4)
}
print(f"SVM Results: {results['banknote']['SVM']} \n")
print(f"Best SVM parameters: {grid_search_svm_banknote.best_params_} \n")
print("\n-----\n")

print("Haberman Dataset \n")

results['haberman'] = {}

# 1. Naive Bayes

print("Training Naive Bayes \n")

nb_haberman = GaussianNB()
nb_haberman.fit(X_haberman_train, y_haberman_train)
y_haberman_pred_nb = nb_haberman.predict(X_haberman_test)
results['haberman']['Naive Bayes'] = {
    'Accuracy': round(accuracy_score(y_haberman_test,
y_haberman_pred_nb), 4),
    'Macro F1-Score': round(f1_score(y_haberman_test,
y_haberman_pred_nb, average='macro'), 4),
    'Macro Precision': round(precision_score(y_haberman_test,
y_haberman_pred_nb, average='macro'), 4),
    'Macro Recall': round(recall_score(y_haberman_test,
y_haberman_pred_nb, average='macro'), 4)
}

print(f"Naive Bayes Results: {results['haberman']['Naive Bayes']} \n")
print("\n-----\n")

# 2. Logistic Regression

print("Training Logistic Regression \n")

lr_haberman = LogisticRegression(random_state=42)
param_grid_lr_haberman = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'solver': ['liblinear']}
grid_search_lr_haberman = GridSearchCV(lr_haberman,
param_grid_lr_haberman, cv=5, scoring='f1_macro')
grid_search_lr_haberman.fit(X_haberman_train, y_haberman_train)
best_lr_haberman = grid_search_lr_haberman.best_estimator_
y_haberman_pred_lr = best_lr_haberman.predict(X_haberman_test)
results['haberman']['Logistic Regression'] = {
    'Accuracy': round(accuracy_score(y_haberman_test,

```

```

y_haberman_pred_lr), 4),
    'Macro F1-Score': round(f1_score(y_haberman_test,
y_haberman_pred_lr, average='macro'), 4),
    'Macro Precision': round(precision_score(y_haberman_test,
y_haberman_pred_lr, average='macro'), 4),
    'Macro Recall': round(recall_score(y_haberman_test,
y_haberman_pred_lr, average='macro'), 4)
}

print(f"Logistic Regression Results: {results['haberman']['Logistic
Regression']} \n")
print(f"Best Logistic Regression parameters:
{grid_search_lr_haberman.best_params_} \n")
print("\n-----\n")

# 3. Support Vector Machine (SVM)

print("Training SVM \n")

svm_haberman = SVC(random_state=42)
param_grid_svm_haberman = {'C': [0.1, 1, 10, 100], 'kernel':
['linear', 'rbf'], 'gamma': ['scale', 'auto']}
grid_search_svm_haberman = GridSearchCV(svm_haberman,
param_grid_svm_haberman, cv=5, scoring='f1_macro')
grid_search_svm_haberman.fit(X_haberman_train, y_haberman_train)
best_svm_haberman = grid_search_svm_haberman.best_estimator_
y_haberman_pred_svm = best_svm_haberman.predict(X_haberman_test)
results['haberman']['SVM'] = {
    'Accuracy': round(accuracy_score(y_haberman_test,
y_haberman_pred_svm), 4),
    'Macro F1-Score': round(f1_score(y_haberman_test,
y_haberman_pred_svm, average='macro'), 4),
    'Macro Precision': round(precision_score(y_haberman_test,
y_haberman_pred_svm, average='macro'), 4),
    'Macro Recall': round(recall_score(y_haberman_test,
y_haberman_pred_svm, average='macro'), 4)
}

print(f"SVM Results: {results['haberman']['SVM']} \n")
print(f"Best SVM parameters: {grid_search_svm_haberman.best_params_} \
n")
print("\n-----\n")

```

Banknote Authentication Dataset

Training Naive Bayes

Naive Bayes Results: {'Accuracy': 0.8582, 'Macro F1-Score': 0.856, 'Macro Precision': 0.8571, 'Macro Recall': 0.8551}

Training Logistic Regression

Logistic Regression Results: {'Accuracy': 0.9855, 'Macro F1-Score': 0.9853, 'Macro Precision': 0.9841, 'Macro Recall': 0.9869}

Best Logistic Regression parameters: {'C': 100, 'solver': 'liblinear'}

Training SVM

SVM Results: {'Accuracy': 1.0, 'Macro F1-Score': 1.0, 'Macro Precision': 1.0, 'Macro Recall': 1.0}

Best SVM parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}

Haberman Dataset

Training Naive Bayes

Naive Bayes Results: {'Accuracy': 0.7581, 'Macro F1-Score': 0.57, 'Macro Precision': 0.686, 'Macro Recall': 0.572}

Training Logistic Regression

Logistic Regression Results: {'Accuracy': 0.7581, 'Macro F1-Score': 0.5338, 'Macro Precision': 0.7147, 'Macro Recall': 0.5516}

Best Logistic Regression parameters: {'C': 0.001, 'solver': 'liblinear'}

Training SVM

SVM Results: {'Accuracy': 0.7258, 'Macro F1-Score': 0.5127, 'Macro Precision': 0.5772, 'Macro Recall': 0.5299}

Best SVM parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}

Thank you