



BITS Pilani
Pilani Campus

Popular scaling approaches

Prof. Akanksha Bharadwaj
CSIS Department



BITS Pilani
Pilani Campus

SE ZG583, Scalable Services

Lecture No. 2



BITS Pilani
Pilani Campus



Partitioning and Sharding

Introduction



- Partitioning and Sharding are scalable approaches to manage and distribute large volumes of data in a database system.
- In many large-scale solutions, data is divided into *partitions* that can be managed and accessed separately.

Why partition data?

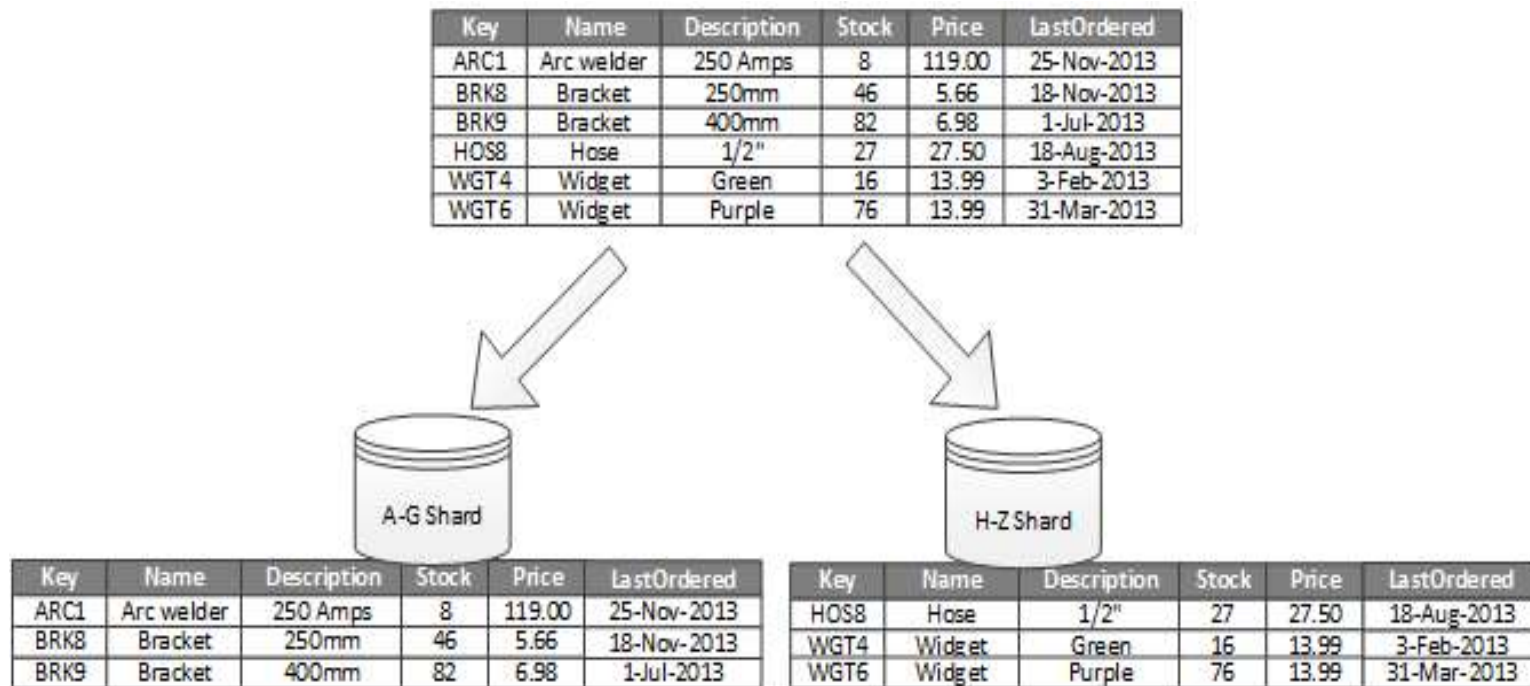
- Improve scalability
- Improve performance
- Improve security
- Provide operational flexibility
- Improve availability

Types of Partitioning




- Horizontal partitioning
- Vertical partitioning
- Functional partitioning

Horizontal partitioning (Sharding)

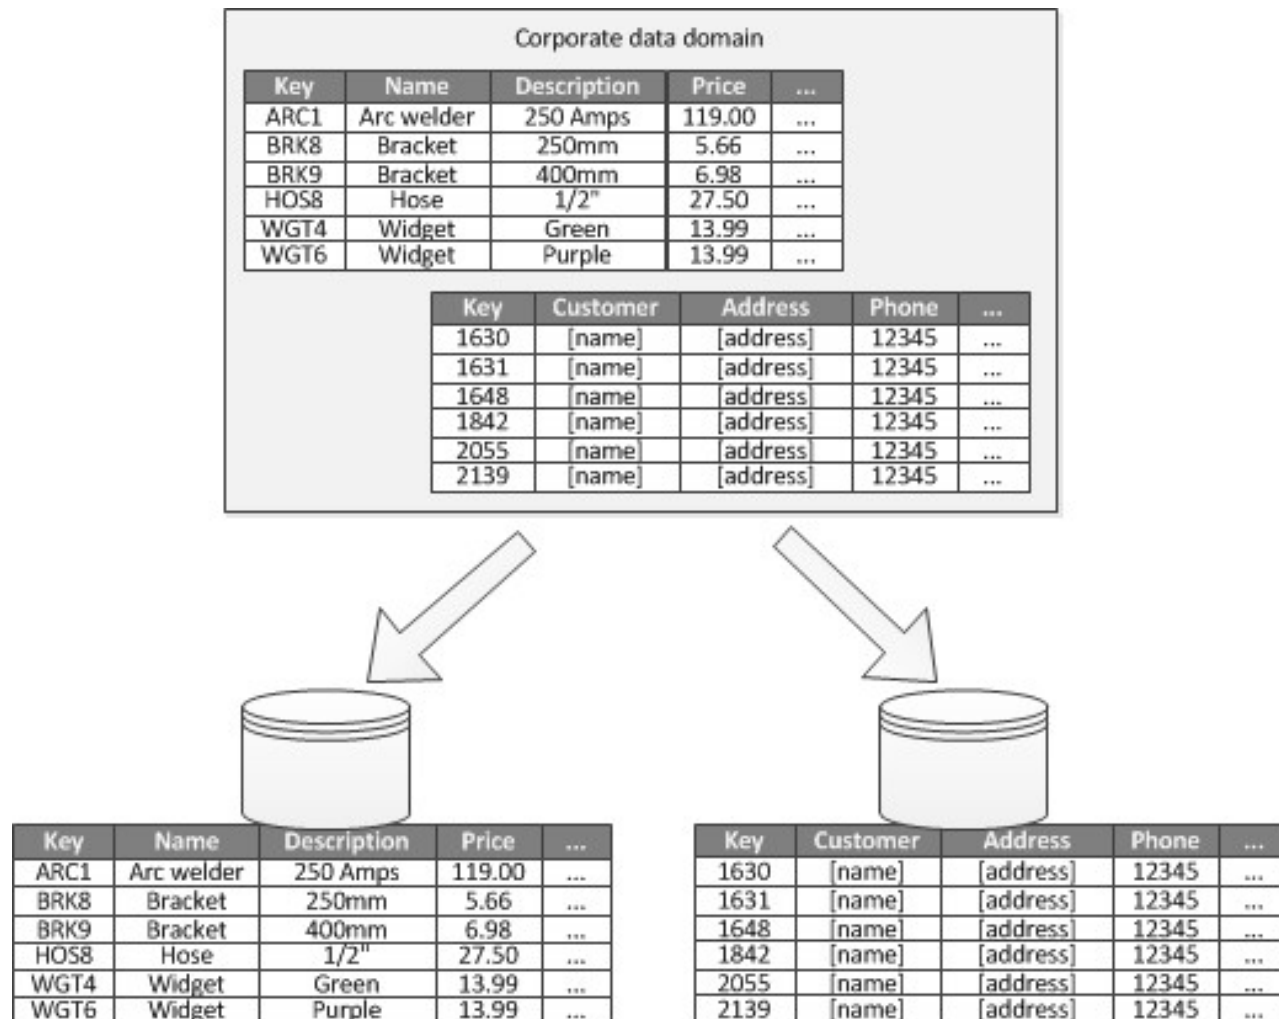


Vertical partitioning

Key	Name	Description	Stock	Price	LastOrdered
ARC1	Arc welder	250 Amps	8	119.00	25-Nov-2013
BRK8	Bracket	250mm	46	5.66	18-Nov-2013
BRK9	Bracket	400mm	82	6.98	1-Jul-2013
HOS8	Hose	1/2"	27	27.50	18-Aug-2013
WGT4	Widget	Green	16	13.99	3-Feb-2013
WGT6	Widget	Purple	76	13.99	31-Mar-2013



Functional partitioning





BITS Pilani
Pilani Campus



Scalability using NoSQL and HDFS

NoSQL and Scalability

NoSQL databases (e.g., MongoDB, Cassandra, DynamoDB) are specifically built for high scalability, especially in scenarios requiring distributed systems, real-time data access, and schema flexibility.

Scalability Features in NoSQL:

- Horizontal Scalability (with Sharding)
- Dynamic Schema
- Replication
- High Throughput
- Decentralized Architecture

Use Cases for NoSQL Scalability



- Social media platforms managing large volumes of user interactions.
- E-commerce systems requiring fast, scalable databases for product catalogs and transactions.
- Real-time analytics on streaming data.

HDFS and Scalability

HDFS is a distributed file system designed for storing and processing massive datasets (big data) in a scalable and fault-tolerant manner. It forms the foundation of the Hadoop ecosystem.

Scalability Features in HDFS:

- Horizontal Scalability
 - Replication for Fault Tolerance
 - Write Once, Read Many (WORM) Design
 - MapReduce Integration
 - Block Storage
 - Cluster Management
-

Use Cases for HDFS Scalability



- Storing and analyzing massive datasets (e.g., petabytes of log data).
- Batch processing for data warehouses and ETL pipelines.
- Data lakes for structured and unstructured data.



BITS Pilani
Pilani Campus



Managing high velocity data streams

Key Challenges



- High Throughput
- Low Latency
- Scalability
- Fault Tolerance
- Storage Efficiency



BITS Pilani
Pilani Campus

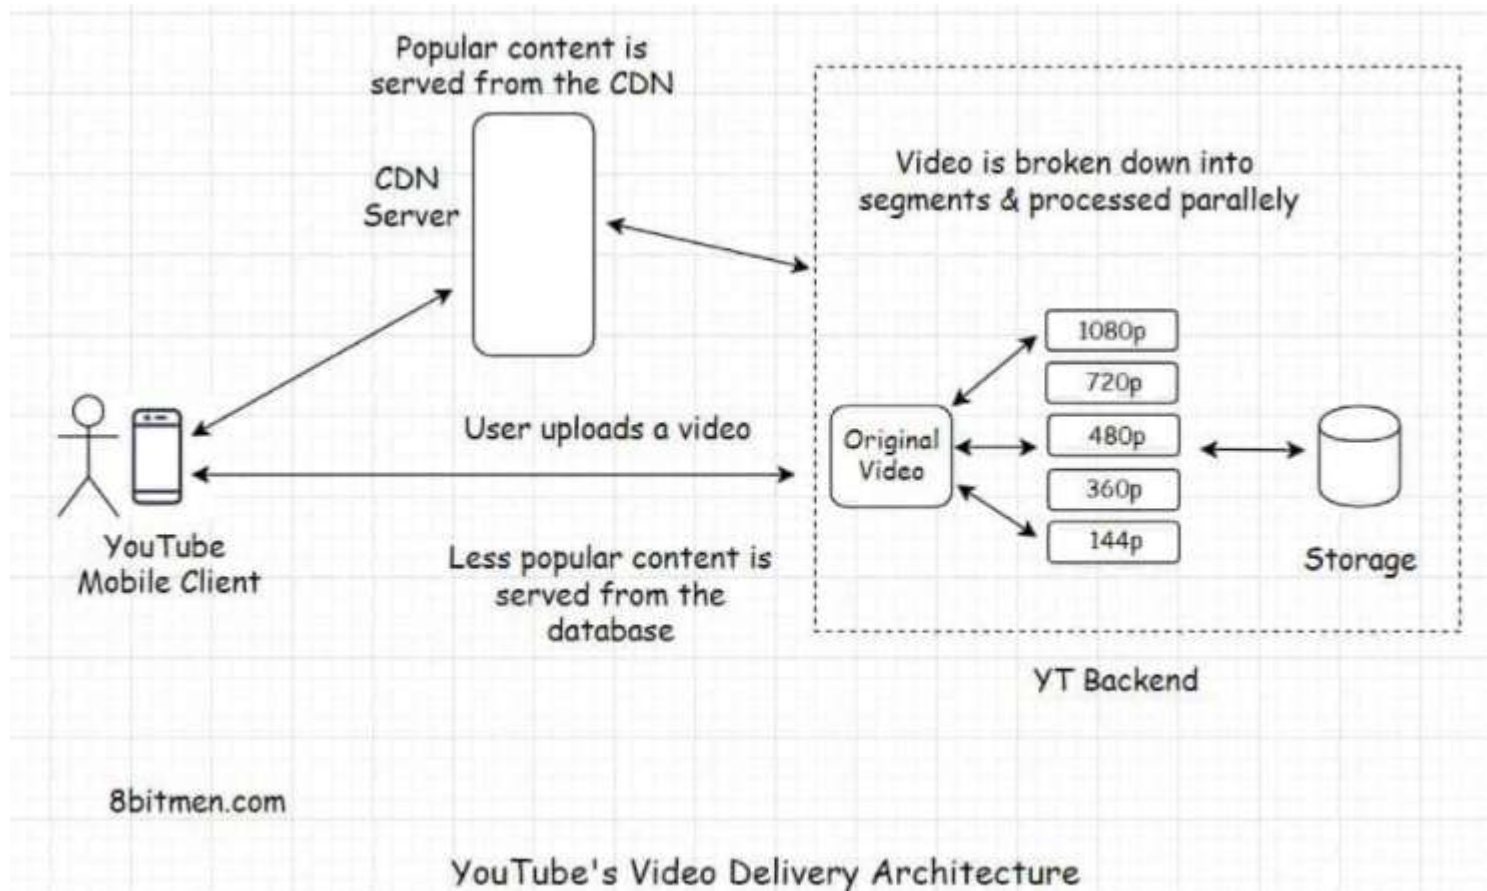
Strategies for Managing High-Velocity Video Data

1. Content Delivery Network



- CDN is nothing more than a bunch of globally distributed computers that are directly connected and move data from one end to another.
- CDN stores multiple copies of video content at edge locations.
- We can use CDNs to cache and distribute video content closer to end-users, reducing latency and bandwidth usage.
- Examples: Akamai, Cloudflare, AWS CloudFront.

CDN Case Study: YouTube



2. Adaptive Bitrate Streaming (ABR)



- Deliver video streams in different quality levels (bitrates) and dynamically switch between them based on the user's network speed and device capabilities.
- Protocols That Support ABR: HLS (HTTP Live Streaming), DASH (Dynamic Adaptive Streaming over HTTP).

3. Efficient Video Compression



- Use advanced codecs to reduce video file size while maintaining quality.
- **Example for video Streaming:** To make the videos viewable on different devices Netflix performs transcoding or encoding which involves finding errors and converting the original video into different formats and resolutions.

4. Edge Computing



- Edge computing is a distributed information technology (IT) architecture in which client data is processed at the periphery of the network, as close to the originating source as possible.
- We can process and cache video content at edge locations to reduce latency and offload computation from the central servers.
- Example: Transcode video streams at the edge for faster delivery.

Case Study: Smart Security Camera



- **Device with Edge Capability:** A smart security camera equipped with an edge AI chip processes video feeds locally.
- **Edge Processing:** The camera runs motion detection algorithms (e.g., using TensorFlow Lite) directly on the device to identify unusual activity, like someone entering the camera's field of view.
- **Action at the Edge:** If motion is detected: Capture the event (e.g., save a short video clip). Send an alert to the homeowner's mobile device (e.g., "Motion detected at the front door"). If no motion is detected, discard the video feed to save bandwidth.
- **Cloud Involvement:** Only the processed, important data (e.g., the detected motion event) is sent to the cloud for storage or additional analysis.

5. Real-Time Data Processing



- Process telemetry data (e.g., playback statistics, user behavior) in real time to optimize streaming.
- Use stream processing frameworks like Apache Flink or Kafka Streams to analyze and react to user interactions instantly.




BITS Pilani
Pilani Campus

Real Time Analytics

Case Study: Real Time Fraud Detection



Architecture



- Mobile phone call metadata is sent from the source system to an Azure Event Hubs instance.
- A Stream Analytics job is started, which receives data via the event hub source.
- The Stream Analytics job runs a predefined query to transform the input stream and analyze it based on a fraudulent-transaction algorithm.
- The Stream Analytics job writes the transformed stream representing detected fraudulent calls to an output sink in Azure Blob storage.

Web Conferencing



- Web conferencing involves real-time audio, video, and data sharing, requiring efficient processing of high-velocity data streams to ensure a seamless user experience.

Steps in Processing Streaming Data for Web Conferencing



1. Data Capture
2. Data Compression
3. Data Transmission
4. Stream Processing at the Server
5. Adaptive Bitrate Streaming
6. Content Delivery
7. Client-Side Rendering
8. Data Analytics and Feedback

Real-World Example: Zoom

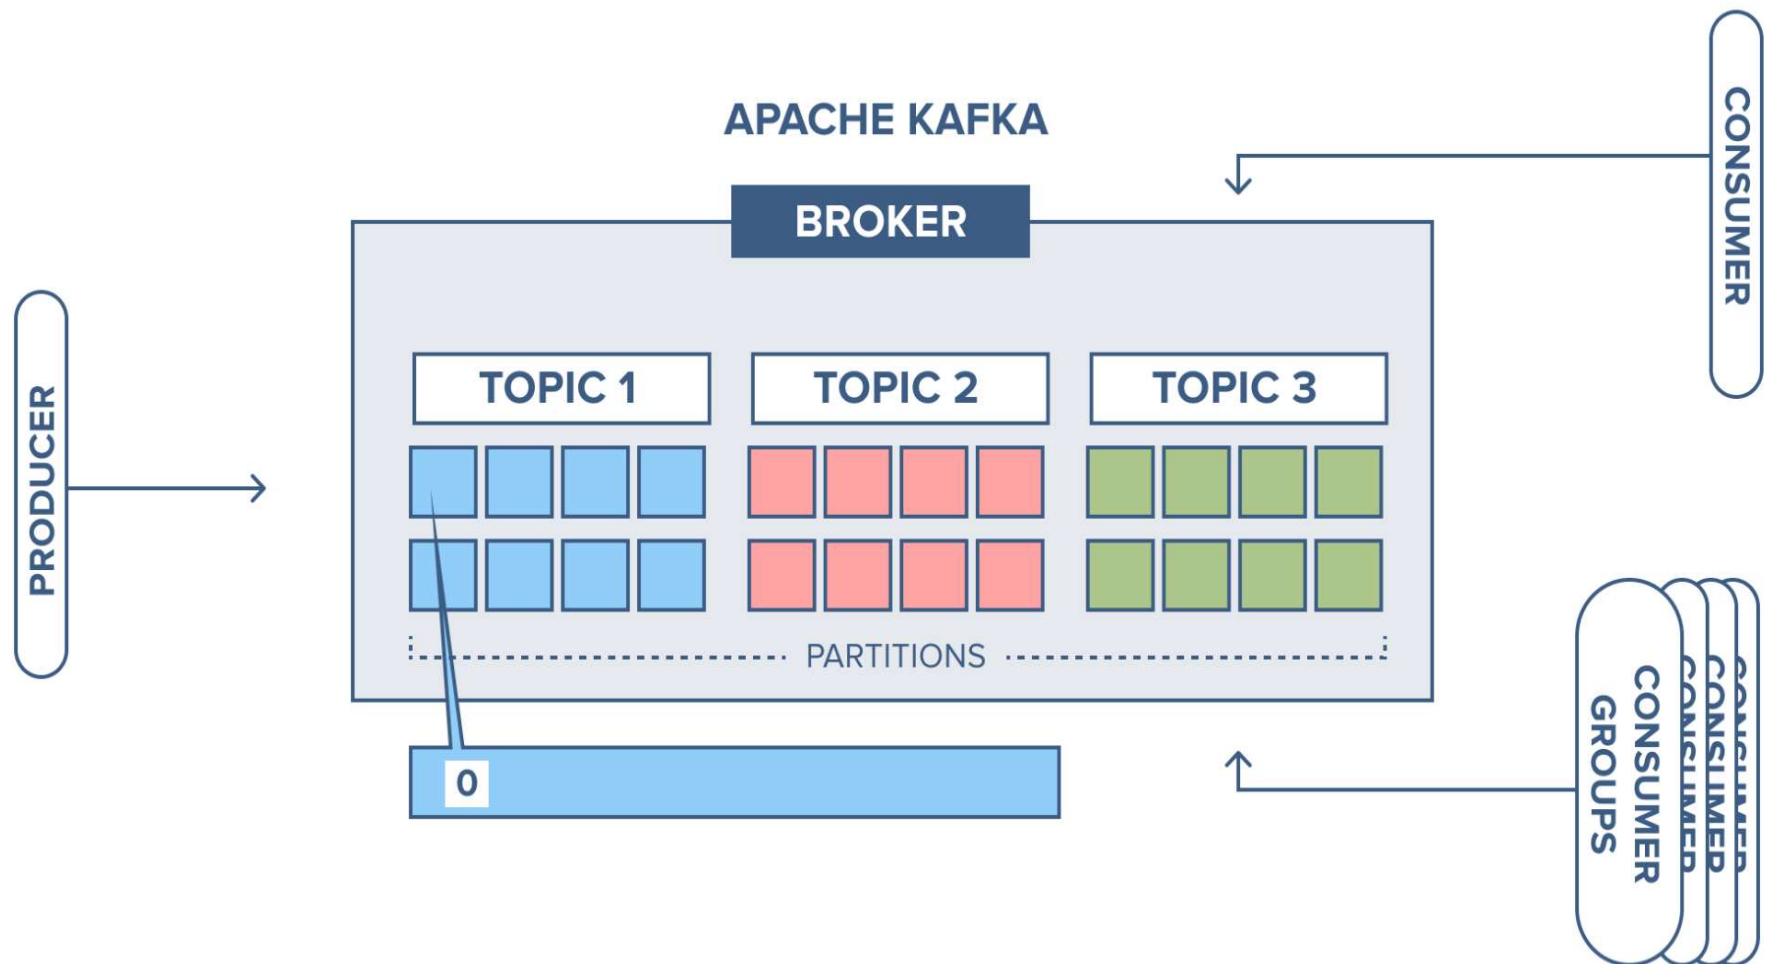


- **Data Capture:** Captures audio, video, and chat data from participants.
- **Compression:** Compresses streams to optimize for bandwidth.
- **Streaming Servers:** Use SFU architecture to minimize latency.
- **Scaling:** Scales dynamically using cloud servers to accommodate user load.

Managing high velocity Data Streams using Kafka



- **Kafka** is a powerful, distributed streaming platform designed to handle massive amounts of data in real-time. It's widely used for building real-time data pipelines and applications.



Kafka for Stream Processing and Analytics



- **Use Case:** Processing and analyzing continuous streams of data for real-time insights, anomaly detection, and predictive analytics.
- **Example:** Financial institutions use Kafka to process market data feeds, detect trading anomalies, and make real-time trading decisions. Retailers analyze customer behavior and preferences in real-time to offer personalized recommendations and promotions.

Ref: <https://medium.com/@psnavya90/>

Self Study



Article: <https://netflixtechblog.com/behind-the-streams-live-at-netflix-part-1-d23f917c2f40>

References



- <https://docs.microsoft.com/en-us/azure/architecture/best-practices/data-partitioning>
- <https://www.mongodb.com/nosql-explained>
- https://hadoop.apache.org/docs/r1.2.1/hdfs_design.html
- <https://netflixtechblog.com/>
- <http://highscalability.com/youtube-architecture>
- <https://docs.microsoft.com/en-us/azure/architecture/example-scenario/data/fraud-detection>
- <https://kafka.apache.org/11/documentation/streams/architecture>
- <https://www.gsma.com/iot/wp-content/uploads/2018/11/loT-Edge-Opportunities-c.pdf>
- <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7696529/>