



**BITS Pilani**  
Pilani Campus

# Popular scaling approaches continued

Prof. Akanksha Bharadwaj  
Asst. Professor, CSIS Department



**BITS Pilani**  
Pilani Campus



# **SE ZG583, Scalable Services**

## **Lecture No. 3**

# Agenda



## Popular scaling approaches

### Managing high volume transactions

- Service Replicas & load balancing
- Minimizing event processing: Command Query Responsibility Segregation (CQRS)
- Asynchronous communication
- Caching techniques: Distributed cache, global cache

### Scalability features in the Cloud (AWS, Azure, Google)

- Auto-scaling
- Horizontal and vertical scaling
- Use of Load balancers
- Virtualization
- Serverless computing
- Best Practices for Achieving Scalability

### Case Study



**BITS Pilani**  
Pilani Campus



# Managing high volume transactions

# Introduction to replication

---

- Replication is the practice of keeping several copies in different places.
- These replicas help distribute the workload and ensure that failures in one instance do not bring down the entire system.
- What is the need of replication?
  - High Availability
  - Fault Tolerance
  - Load sharing

# Service Replicas



- It is a copy of the service logic that runs on one or more nodes of the cluster.
- Types of replicas: These replicas can be categorized based on their deployment model, scaling approach, and management strategy.
  - Stateful
  - Stateless
  - Active-Passive
  - Active-active
  - Leader-follower

# Stateful and Stateless Replicas



- **Stateful Replicas:** Each replica maintains its own persistent state, requiring synchronization across instances.
- **Use Case:** Databases, caching services, distributed file systems.
- **Stateless Replicas:** Each replica operates independently and does not maintain any session-specific data.
- **Use Case:** API gateways, web servers, RESTful microservices.

# Active-Active and Active-Passive Replicas



- **Active-Active:** All replicas handle requests simultaneously in a distributed load-sharing model.
- **Use Case:** High-traffic applications needing high availability.
- **Active-Passive:** One primary instance handles traffic, while replicas remain on standby and become active in case of failure.
- **Use Case:** Database replication, disaster recovery systems.



# Leader-Follower Replica

---

- A replication model where a **Leader (Primary)** node handles **write operations** and **Follower (Replica)** nodes handle **read operations**. If the Leader fails, a Follower is promoted as the new Leader.
- Usecases: Databases, Caching system, Message Brokers

# Load Balancing



- **Load balancing** refers to efficiently distributing incoming network traffic across a group of backend servers, also known as a *server farm* or *server pool*.
- If a single server goes down, the load balancer redirects traffic to the remaining online servers.
- When a new server is added to the server group, the load balancer automatically starts to send requests to it.

# What are load balancers?



- A load balancer is a software or hardware device that keeps any one server from becoming overloaded.
- A load-balancing algorithm is the logic that a load balancer uses to distribute network traffic between servers
- Types Based on Deployment Location:
  - Software-based load balancers
  - Hardware-based load balancers
  - Cloud Load Balancer

# Why is Load Balancing Important for Scalability?

---



Load balancing ensures:

- Even distribution of requests across multiple instances.
- Reduced latency and improved response time by directing traffic to less loaded servers.
- Fault tolerance and high availability by rerouting traffic in case of server failures.
- Efficient resource utilization and better cost management.

# Command Query Responsibility Segregation (CQRS)



- It is a pattern that separates read and update operations for a data store.
- This pattern is particularly useful in **highly scalable and distributed applications**.

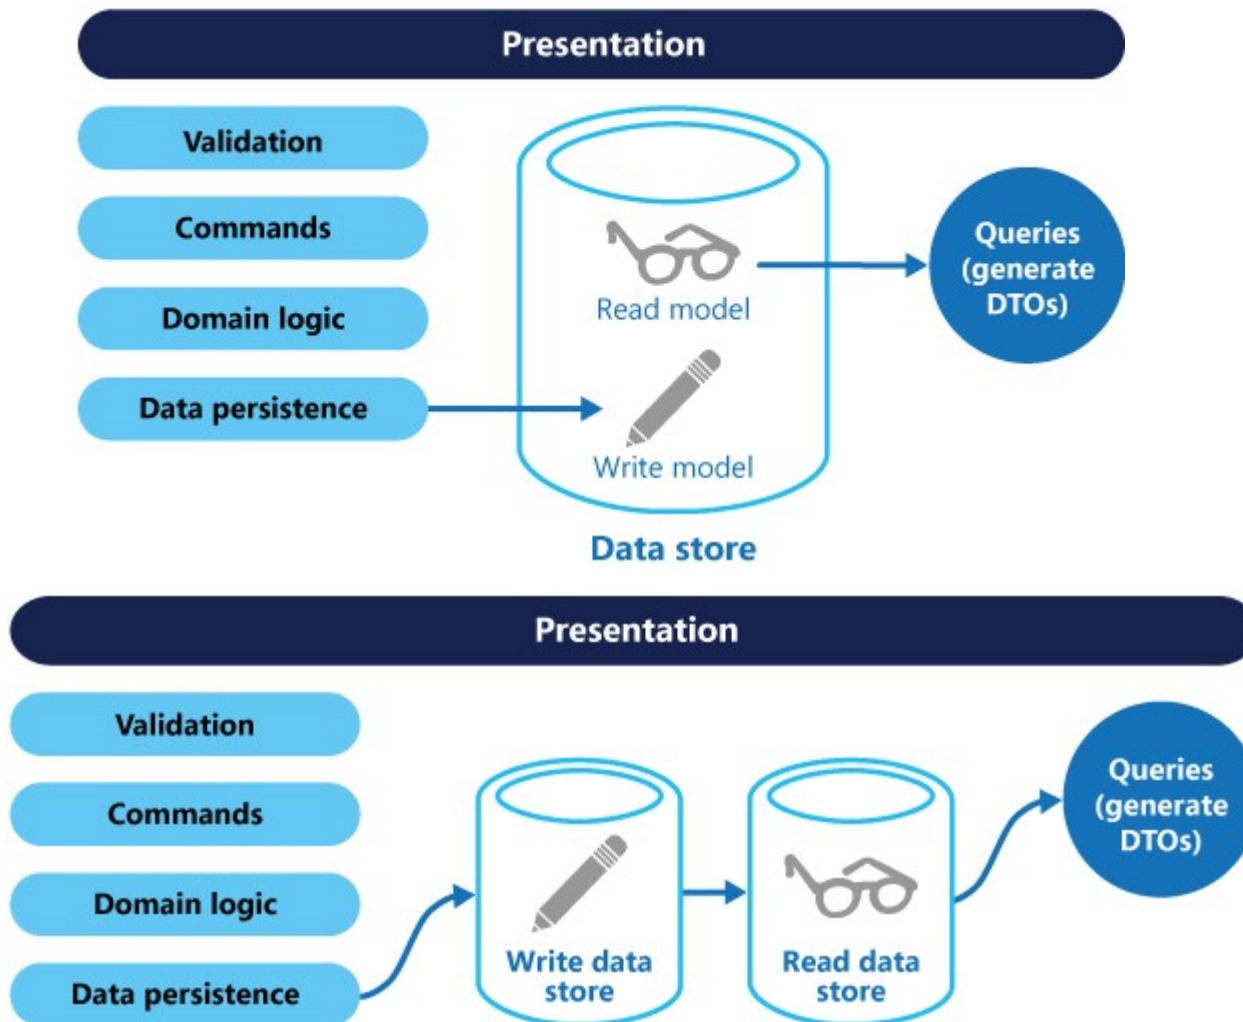
# Why CQRS for Scalability?



CQRS improves **scalability** by:

- Separating concerns
- Independent scaling
- Optimizing for reads and writes separately
- Supporting multiple databases
- Simpler queries

# How CQRS works?



# Some challenges of implementing CQRS

---



- Complexity
- Messaging
- Eventual consistency



# What is Event sourcing?



- **Event Sourcing** is an architectural pattern where **state changes** in a system are stored as a **sequence of events** instead of modifying the current state directly.

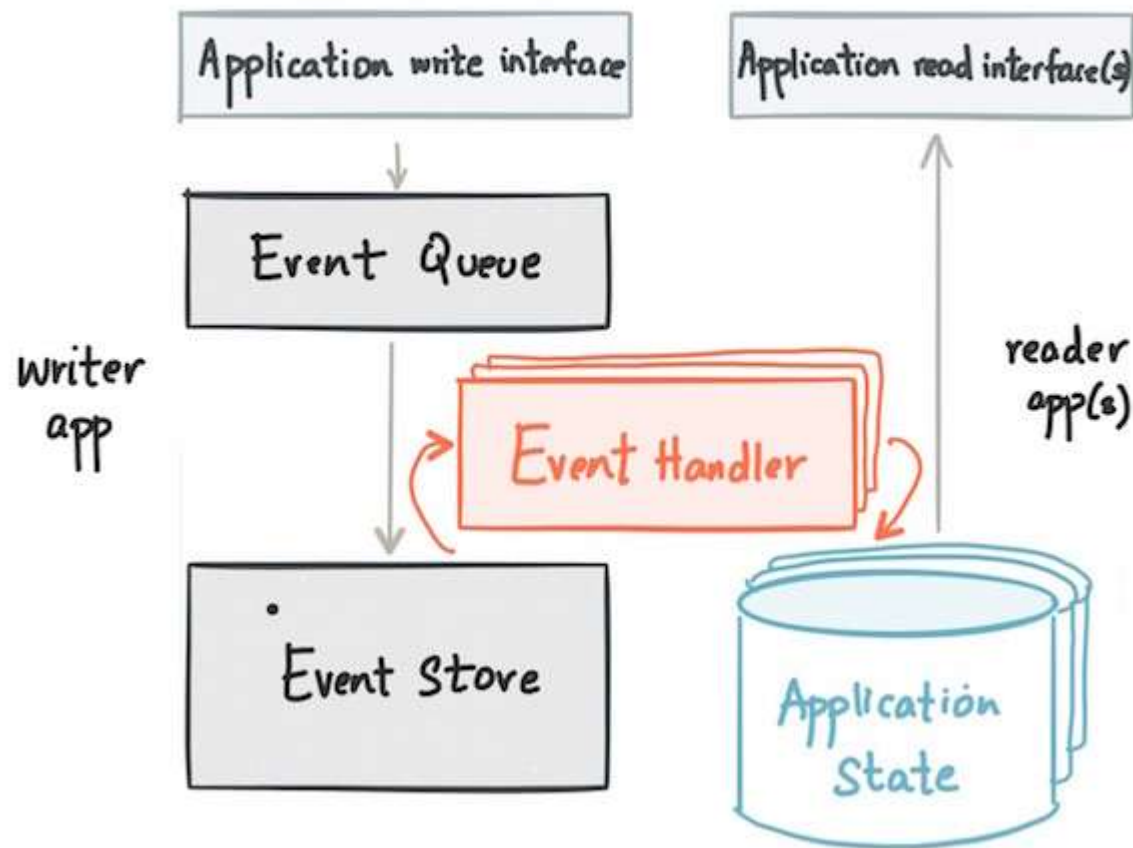
# How Event Sourcing and CQRS Work Together

---



- **Write operations (Commands)** → Generate **immutable events** instead of updating the database directly.
- **Events are stored** in an **event store** (e.g., Kafka, EventStoreDB).
- **Read operations (Queries)** → Use a **separate, read-optimized database** that is updated asynchronously.

# Event Sourcing and CQRS



<https://mrwersa.medium.com/cqrs-pattern-with-kafka-streams-part-1-112f381e9b98>

# Architecture of CQRS with Event Sourcing



## A. Command Side (Write Model)

- Handles create, update, delete operations.
- Instead of updating a relational DB, writes generate events (e.g., OrderPlaced).
- Events are stored in an event store and published to event consumers.

## B. Event Store

- A log of all past events, acting as the source of truth.
- Example storage: Kafka, EventStoreDB, PostgreSQL (JSONB), DynamoDB Streams.



## C. Query Side (Read Model)

- Events from the event store are processed to update a read-optimized database.
- Read models can use SQL (denormalized tables), NoSQL (MongoDB, Redis), or search engines (Elasticsearch).

## D. Event Handlers & Projections

- As events are stored, event handlers process them to update projections (optimized views for queries).
- Example: A UserRegistered event updates a UserProfile table in a read database

# Protocols for communication



In **distributed systems**, communication between components plays a crucial role in **scalability, performance, and reliability**.

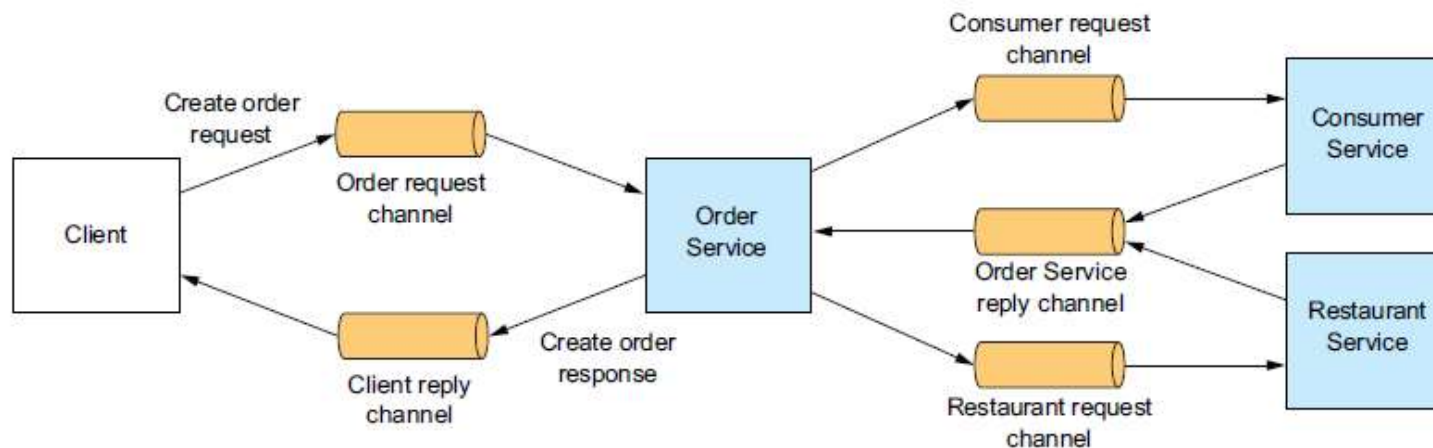
Two main communication patterns are:

- **Synchronous Communication:** The sender waits for a response before proceeding.
- **Asynchronous Communication:** The sender does not wait for an immediate response, improving system efficiency.

# Asynchronous Communication



- Services communicating by exchanging messages over messaging channels.



# Synchronous Vs Asynchronous



Feature	Synchronous (Blocking)	Asynchronous (Non-Blocking)
Use Case	Request-response APIs	Event-driven systems, messaging
Latency	Higher due to waiting	Lower, as operations continue
Scalability	Limited by concurrent requests	High scalability
Resilience	Prone to failures/timeouts	More resilient to failures
Complexity	Easier to implement	More complex (requires event handling)
Examples	REST APIs, RPC calls	Kafka, RabbitMQ, WebSockets



# Message Broker

---

- It is a way of implementing asynchronous communication
- A message broker is an intermediary through which all messages flow.

Examples of popular open source message brokers include the following:

- ActiveMQ
- RabbitMQ
- Apache Kafka

# Benefits of Message Broker

---



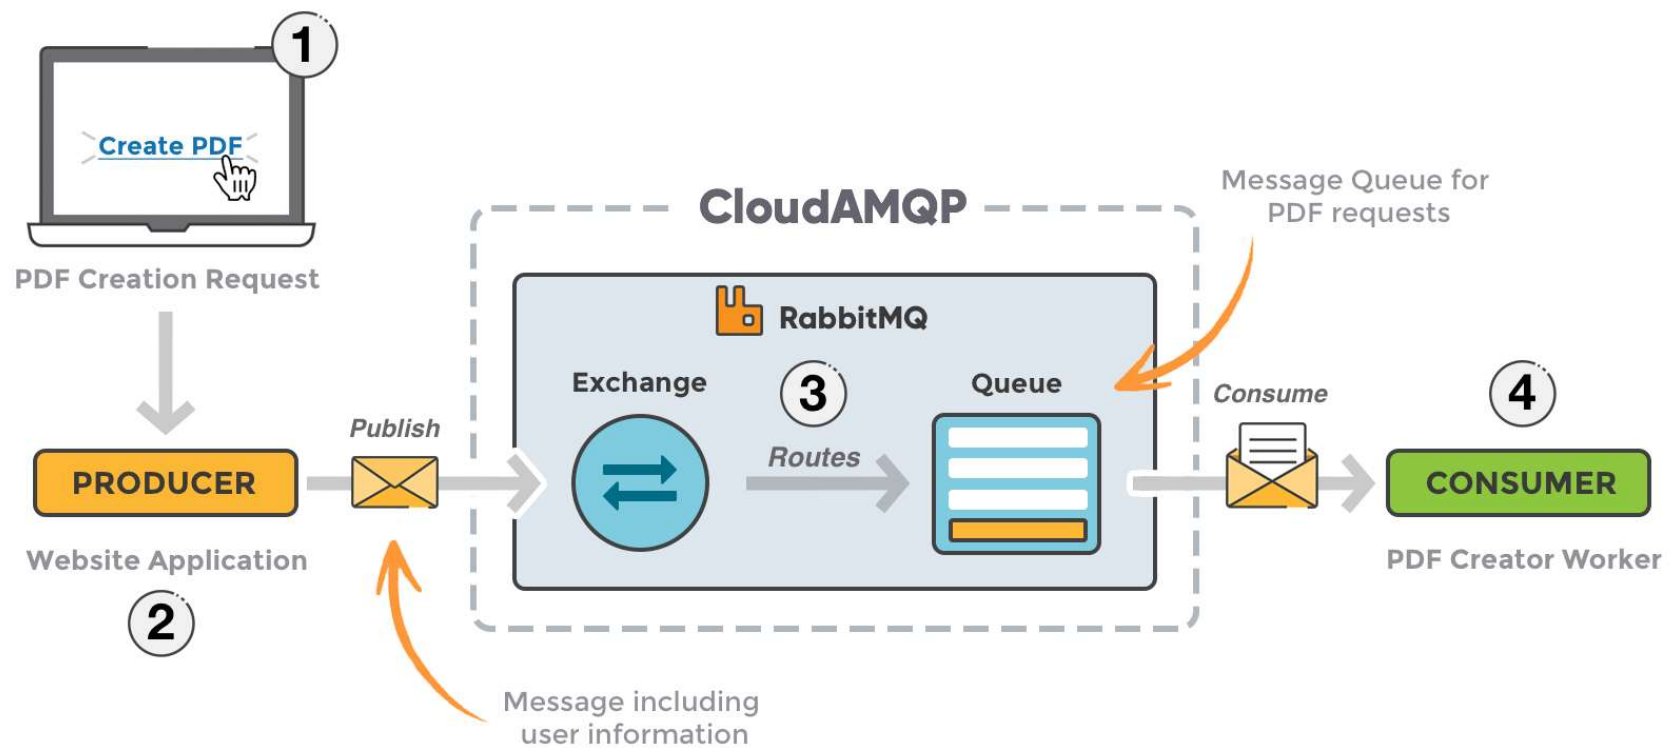
- *Loose coupling*
- *Message buffering*
- *Explicit interprocess communication*
- *Resiliency*

# Drawbacks of Message Broker



- *Potential performance bottleneck*
- *Potential single point of failure*
- *Additional operational complexity*

# Example



<https://www.cloudamqp.com/blog/part1-rabbitmq-for-beginners-what-is-rabbitmq.html>

# What is Caching?



- Caching is a critical technique in **distributed systems** to **reduce latency, improve performance, and optimize resource usage**.
- Caching allows you to efficiently reuse previously retrieved or computed data.

# Types of cache




Different types of caching strategies are used depending on system architecture, data consistency needs, and workload patterns.

- Server side
- Client side
- Application-Level Cache
- Distributed Cache

# Distributed Caches

- A distributed cache may span multiple servers so that it can grow in size and in transactional capacity.
- It is mainly used to store application data residing in database and web session data.



# Advantages of Distributed Cache

---



When cached data is distributed, the data:


- Is *coherent* (consistent) across requests to multiple servers.
- Survives server restarts and app deployments.
- Doesn't use local memory.



# Example: Memcached




## Without Memcached



When Used Separately  
Total Usable Cache size: **64MB**

## With Memcached



When Logically Combined  
Total Usable Cache size: **128MB**

# Global Caches



- **Global Cache** refers to a caching mechanism that operates across multiple geographic regions or data centers, ensuring faster data access and improved system scalability.

## **Key Characteristics of Global Cache:**

- ✓ Distributed Across Multiple Regions
- ✓ Reduces Latency
- ✓ Load Balancing
- ✓ High Availability

# Google Global Cache (GGC)



- **Google Global Cache (GGC)** is a specialized caching system deployed by **Google in Internet Service Providers' (ISPs) networks** to improve the delivery of Google services such as:
  - YouTube
  - Google Search
  - Google Drive
  - Google Play Store
- GGC is part of **Google's Content Delivery Network (CDN)**, designed to **reduce bandwidth costs and improve access speed** by caching content **closer to users**.



**BITS Pilani**  
Pilani Campus



# Scalability features in the Cloud

# Auto-scaling



- Auto Scaling is a technique used in **cloud computing and distributed systems** to **dynamically adjust computing resources** based on real-time demand.

# Types of Auto Scaling



- **Vertical Scaling** is defined as increasing a single machine's capacity with the rising resources in the same logical server or unit. .
- **Horizontal Scaling** is an approach to enhance the performance of the server node by adding new instances of the server to the existing servers to distribute the workload equally



# Comparison



Vertical Vs Horizontal Scaling	Vertical scaling	Horizontal Scaling
<b>Data</b>	Data is executed on a single node	Data is partitioned and executed on multiple nodes
<b>Data Management</b>	Easy to manage – share data reference	Complex task as there is no shared address space
<b>Downtime</b>	Downtime while upgrading the machine	No downtime
<b>Upper limit</b>	Limited by machine specifications	Not limited by machine specifications
<b>Cost</b>	Lower licensing fee	Higher licensing fee

Image: Google

# What is virtualization?



- Virtualization is a technology that allows multiple **virtual instances** (e.g., virtual machines, containers) to run on a single **physical machine** by abstracting hardware resources.



# Types of virtualization



- Server virtualization
- Storage virtualization
- Data virtualization
- Application virtualization
- Containerization

.....

# How Virtualization Improves Scalability?

---



- On-Demand Resource Allocation
- Efficient Load Balancing
- Multi-Tenancy Support
- High Availability & Fault Tolerance
- Cost-Effective Scaling

# What is Serverless Computing?

---



- Serverless computing is a cloud computing model where cloud providers automatically manage the infrastructure and dynamically allocate resources as needed.

# Key Characteristics of Serverless Computing

---



- No Server Management
- Event-Driven
- Automatic Scaling
- Pay-as-You-Go

# How Serverless Computing Enhances Scalability?

---



- Auto Scaling Based on Demand
- Instant Scaling and Elasticity
- Cost Efficiency
- Statelessness
- Managed Load Balancing

# Types of Serverless Architectures for Scalability

---



- **Function-as-a-Service (FaaS):** developers write individual functions that are executed in response to specific events
- **Backend-as-a-Service (BaaS):** use of managed backend services that are provided by cloud vendors

# Challenges of Serverless Computing for Scalability

---



- ✗ Cold Start Latency
- ✗ State Management
- ✗ Vendor Lock-In
- ✗ Debugging and Monitoring

# Best Practices for Achieving Scalability

---



- Be Stateless
- Use Load Balancers
- Auto Scaling
- Optimize Database Performance
- Use Caching Mechanisms
- Microservices Architecture
- Event-Driven Architecture
- Use Serverless Computing
- Monitor and Analyze Performance

By following these best practices, you can build highly scalable applications that handle increased demand without compromising performance or reliability.



# Case Study: Amazon Prime



- Amazon Prime Video Uses AWS to Deliver Solid Streaming Experience to More Than 18 Million Football Fans
- Amazon Prime needed an architecture that could quickly scale, handle spikes, and have sufficient caching in different layers to manage the demand.
  - Scalable Cloud Infrastructure
  - Content Delivery with Amazon CloudFront
  - High-Performance Video Encoding and Storage
  - Serverless Technology
  - Real-Time Analytics with Amazon Kinesis
  - Global Availability and Reliability
  - Enhanced Security and Compliance

# Self Study



- Amazon Prime case study: <https://aws.amazon.com/solutions/case-studies/amazon-prime-video/>
- Hotstar case study: <https://laveena-j-21.medium.com/cloud-computing-aws-and-disney-hotstar-case-study-f4a3be4669a>
- Article on serverless: <https://www.simform.com/blog/serverless-architecture-guide/>

# References



- [Book](#): The Art of Scalability: Scalable Web Architecture, Processes, and Organizations for the Modern Enterprise, Second Edition by Michael T. Fisher; Martin L. Abbott Published by Addison-Wesley Professional, 2015
- <https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-concepts-replica-lifecycle>
- <https://www.citrix.com/en-in/solutions/application-delivery-controller/load-balancing/what-is-load-balancing.html#:~:text=Load%20balancing%20is%20defined%20as,server%20capable%20of%20fulfilling%20them.>
- <https://docs.microsoft.com/en-us/azure/architecture/patterns/cqrs#:~:text=CQRS%20stands%20for%20Command%20and,performance%2C%20scalability%2C%20and%20security.>
- <https://docs.microsoft.com/en-us/dotnet/architecture/microservices/architect-microservice-container-applications/communication-in-microservice-architecture>
- <https://memcached.org/>
- <https://www.ibm.com/docs/en/integration-bus/10.0?topic=caching-data-overview>
- <https://support.google.com/interconnect/answer/9058809?hl=en>
- <https://aws.amazon.com/>
- <https://www.nginx.com/resources/glossary/load-balancing/>
- <https://www.ibm.com/in-en/cloud/learn/virtualization-a-complete-guide>
- <https://www.confluent.io/blog/event-sourcing-cqrs-stream-processing-apache-kafka-whats-connection/>
- <https://aws.amazon.com/solutions/case-studies/amazon-prime-video/>